

Online Learning Algorithms

Nicolò Cesa-Bianchi¹ and Francesco Orabona²

¹Department of Computer Science & DSRC, Università degli Studi di Milano, Milano, Italy, 20133; email: nicolo.cesa-bianchi@unimi.it

²Department of Electrical and Computer Engineering, Boston University, Boston, MA, USA, 02215; email: francesco@orabona.com

Xxxx. Xxx. Xxx. Xxx. YYYY. AA:1–27

[https://doi.org/10.1146/\(\(please add article doi\)\)](https://doi.org/10.1146/((please add article doi)))

Copyright © YYYY by Annual Reviews.
All rights reserved

Keywords

regret minimization, convex optimization, regression, classification

Abstract

Online learning is a framework for the design and analysis of algorithms that build predictive models by processing data one at the time. Besides being computationally efficient, online algorithms enjoy theoretical performance guarantees that do not rely on statistical assumptions on the data source. In this survey, we describe some of the most important algorithmic ideas behind online learning and explain the main mathematical tools for their analysis. Our reference framework is online convex optimization, a sequential version of convex optimization within which most online algorithms are formulated. More specifically, we provide an in-depth description of Online Mirror Descent and Follow the Regularized Leader, two of the most fundamental algorithms in online learning. As the tuning of parameters is a typically difficult task in sequential data analysis, in the last part of the survey we focus on coin-betting, an information-theoretic approach to the design of parameter-free online algorithms with good theoretical guarantees.

Contents

1. INTRODUCTION	2
1.1. Online convex optimization	3
1.2. Regret and sequential risk	4
1.3. Lower bounds	4
2. ONLINE MIRROR DESCENT	5
2.1. The AdaGrad algorithm	10
3. FOLLOW THE REGULARIZED LEADER	10
3.1. Online Newton Step.....	13
3.2. Online linear regression	15
3.3. Optimistic updates	16
4. UNCONSTRAINED ONLINE CONVEX OPTIMIZATION.....	17
5. OTHER NOTIONS OF REGRET	21

1. INTRODUCTION

The growing success of technologies based on machine learning is driven by the availability of massive data sets in digital format. Processing these large amounts of data poses computational challenges that are not always properly addressed by traditional statistical learning methods. For this reason, online or sequential learning, a framework specifically designed to cope with big data scenarios, has become a key tool in machine learning applications. Online algorithms go through the data points sequentially, using each new data point to adjust their predictive model or estimator. Typically, this adjustment is “local”, as it only involves the current model and the new data point. As each data point is often processed in constant time, this results in an overall running time scaling linearly with the number of data points. Besides the computational advantage, there are more reasons for which sequential learning may be preferred over other approaches. In many application domains—such as online advertising, digital markets, sensor networks, mobile user applications—new data are generated at high rates. In these cases, the sequential adaptation process of online learning has the potential of capturing subtle nonstationary features of the unknown data source.

A fundamental issue in machine learning is what mathematical assumptions on data sources are reasonable to make. Online learning advocates an approach in which the source is viewed as an arbitrary and unknown *deterministic* process. This is a radical departure from classical statistical approaches to sequential decision-making—such as Bayesian decision theory (Berger 2013) or Markov decision processes (Puterman 2014)—and finds its roots in the pioneering works on repeated games by Robbins (1951), Hannan (1957), and Blackwell (1956), where the data source consists of the opponent’s plays in a two-person game. The theme of predicting individual, deterministic sequences also surfaced in other disciplines, including information theory (Cover 1965, Feder et al. 1992), and computer science (Borodin & El-Yaniv 2005). More recently, some of the online learning techniques—like exponential weighted aggregation—have also appeared in the statistical literature (see, e.g., Dalalyan & Tsybakov 2008, Dalalyan et al. 2012, Rigollet et al. 2012).

Stripping the data source of any statistical assumption allows to define a crisp, minimalist framework for investigating the notion of algorithmic learning, where only the

empirical properties of the observed data sequence matter. A substitute for the notion of statistical risk must then be introduced to define a notion of minimax optimality over the sequence of losses in a mathematically rigorous way. In view of that, one should note that in online learning we never measure the predictive power of a single model, but rather consider the ensemble of models sequentially generated by the online learner while processing a data sequence. The notion of risk that we use to evaluate this ensemble is appropriately called *sequential risk*, and measures the extent to which each model generated by the algorithm is able to predict the next element in the sequence. Sequential risk is thus associated with the behavior of an algorithm on an individual data sequence. Based on sequential risk, we then derive the notion of *regret*, which can be viewed as the online counterpart of the statistical excess risk measured with respect to a class of predictive models. The control of regret is the main goal in the analysis of online learning algorithms.

In the rest of this survey, we introduce and describe some of the most fundamental online learning algorithms. Our goal is to explain the behavior of these algorithms through the analysis of their regret. Therefore, rather than going through as many as possible of the existing approaches to online learning, we prefer to focus on the conceptual foundations and the main proof techniques. We believe this is a more effective way to keeping alive the interest of someone who wants to know more about this exciting field of research.

1.1. Online convex optimization

The standard framework for the study of parametric online learning with convex losses¹ is known as Online Convex Optimization, or OCO (see, e.g., Shalev-Shwartz 2012, Hazan 2016, McMahan 2017, Orabona 2019). While it is also possible to design and study nonparametric online learning algorithms, (e.g., Hazan & Megiddo 2007, De Rosa et al. 2015, Kuzborskij & Cesa-Bianchi 2017), we focus here on the more common parametric setting.

As we said earlier, the data-generating mechanism of online learning is an unknown and deterministic process. In OCO, the data process is replaced by a deterministic sequence of unknown and convex loss functions ℓ_t , evaluating the performance of the models incrementally generated by the algorithm. For example, if we want to cast linear regression in the OCO framework, then the loss functions ℓ_t take the form $\ell_t(\mathbf{w}) = (\mathbf{w}^\top \mathbf{x}_t - y_t)^2$, where $\mathbf{w} \in \mathbb{R}^d$ is a linear prediction model and $(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots \in \mathbb{R}^d \times \mathbb{R}$ is the underlying deterministic data sequence.

Throughout the paper, the *model space* \mathbb{X} is a convex, closed, and nonempty subset of \mathbb{R}^d , while a *loss function* is any nonnegative and convex function $\ell_t : \mathbb{X} \rightarrow \mathbb{R}$. For any sequence ℓ_1, ℓ_2, \dots of loss functions, an *online learner* A is a sequence A_1, A_2, \dots of mappings with range \mathbb{X} . The learner's model at time t is $\mathbf{w}_t = A_{t-1}(\ell_1, \dots, \ell_{t-1})$, where $\mathbf{w}_1 = A_1 \in \mathbb{X}$ is the default initial model, and the learner's loss at time t is $\ell_t(\mathbf{w}_t)$. While the notation $\mathbf{w}_t = A_{t-1}(\ell_1, \dots, \ell_{t-1})$ highlights the fact that \mathbf{w}_{t+1} can depend on all the past losses, we are especially interested in cases where the update $\mathbf{w}_t \rightarrow \mathbf{w}_{t+1}$ can be done efficiently based only on \mathbf{w}_t and ℓ_t . Also, for certain applications it makes sense to consider special cases of this framework, where additional assumptions besides convexity are made on the loss functions.

¹Partial extensions of online learning to nonconvex losses have been recently considered by Agarwal et al. (2019).

1.2. Regret and sequential risk

The performance of the learner is measured according to the *regret*

$$R_T(\mathbf{u}) = \sum_{t=1}^T (\ell_t(\mathbf{w}_t) - \ell_t(\mathbf{u})) \quad \text{for } \mathbf{u} \in \mathbb{X}$$

Online learning is concerned with the design of algorithms for which $R_T(\mathbf{u})$ grows sublinearly in T for all $\mathbf{u} \in \mathbb{X}$ and irrespective to the loss sequence (the so-called *no-regret* property). The quantity $\frac{1}{T} \sum_{t=1}^T \ell_t(\mathbf{w}_t)$ is sometimes called *sequential risk* (as opposed to the classical statistical risk), thus sublinear regret implies that the excess sequential risk

$$\frac{1}{T} \sum_{t=1}^T \ell_t(\mathbf{w}_t) - \frac{1}{T} \sum_{t=1}^T \ell_t(\mathbf{u})$$

converges to zero for any $\mathbf{u} \in \mathbb{X}$.

Algorithms that enjoy the no-regret property can also be used to solve convex optimization problems $\min_{\mathbf{w} \in \mathbb{X}} f(\mathbf{w})$, which are viewed as instances of OCO with $\ell_t = f$ for all t . Indeed, Jensen's inequality shows that

$$f(\bar{\mathbf{w}}) - \min_{\mathbf{w} \in \mathbb{X}} f(\mathbf{w}) \leq \frac{1}{T} \left(\sum_{t=1}^T \ell_t(\mathbf{w}_t) - \min_{\mathbf{w} \in \mathbb{X}} \sum_{t=1}^T \ell_t(\mathbf{w}) \right) \leq \frac{R_T(\mathbf{w}^*)}{T}$$

where $\bar{\mathbf{w}} = \frac{1}{T}(\mathbf{w}_1 + \dots + \mathbf{w}_T)$ is the average of the iterates \mathbf{w}_t and \mathbf{w}^* is the minimizer of f in \mathbb{X} . Similarly, we may also consider stochastic optimization problems $\min_{\mathbf{w} \in \mathbb{X}} \mathbb{E}[F(\mathbf{w}, \omega)]$, where the random variable $F(\mathbf{w}, \cdot)$ is the stochastic objective and $F(\cdot, \omega)$ is convex for all $\omega \in \Omega$. Given access to i.i.d. draws $\omega_1, \omega_2, \dots$, we can solve stochastic optimization problems using a no-regret algorithm run with $\ell_t(\mathbf{w}) = F(\mathbf{w}, \omega_t)$. Let $\mathbf{w}^* \in \operatorname{argmin}_{\mathbf{w} \in \mathbb{X}} \mathbb{E}[F(\mathbf{w}, \omega)]$ and observe that, using Jensen's inequality once more,

$$\begin{aligned} \mathbb{E}[F(\bar{\mathbf{w}}, \omega) - F(\mathbf{w}^*, \omega)] &\leq \mathbb{E} \left[\frac{1}{T} \sum_{t=1}^T (F(\mathbf{w}_t, \omega_t) - F(\mathbf{w}^*, \omega_t)) \right] \\ &= \mathbb{E} \left[\frac{1}{T} \sum_{t=1}^T (\ell_t(\mathbf{w}_t) - \ell_t(\mathbf{w}^*)) \right] \leq \mathbb{E} \left[\frac{R_T(\mathbf{w}^*)}{T} \right] \end{aligned}$$

In machine learning, stochastic optimization is typically used to solve empirical risk minimization problems

$$\min_{\mathbf{w} \in \mathbb{X}} \frac{1}{m} \sum_{i=1}^m F(\mathbf{w}, \mathbf{z}_i) \tag{1}$$

where $\mathbf{z}_1, \dots, \mathbf{z}_m$ is a dataset and $F(\mathbf{w}, \mathbf{z}_i)$ measures the loss of \mathbf{w} on the data point \mathbf{z}_i . If $F(\cdot, \mathbf{z})$ is convex for all \mathbf{z} , then we may set $\ell_t = F(\cdot, \mathbf{Z}_t)$, where $\mathbf{Z}_1, \mathbf{Z}_2, \dots$ are i.i.d. uniform draws from the dataset.

1.3. Lower bounds

Consider the easy case when \mathbb{X} is a bounded² set with diameter D and all losses ℓ_t are Lipschitz on \mathbb{X} , how well can we control the regret R_T in this scenario? It turns out that

²This can be a plausible assumption in many practical cases. For instance, in online linear regression when upper bounds on $\max_t \|\mathbf{x}_t\|_2$ and $\max_t |y_t|$ are known in advance.

the worst case for OCO occurs when all loss functions are linear. Interestingly, the proof uses a stochastic rather than deterministic loss process. More specifically, let $\mathbf{v}_1, \mathbf{v}_2 \in \mathbb{X}$ such that $\|\mathbf{v}_1 - \mathbf{v}_2\|_2 = D$ and define $\mathbf{z}_0 = \frac{\mathbf{v}_1 - \mathbf{v}_2}{\|\mathbf{v}_1 - \mathbf{v}_2\|_2}$. Stochastic losses L_1, L_2, \dots are defined by $L_t(\mathbf{w}) = \varepsilon_t L \mathbf{z}_0^\top \mathbf{w}$, where $\varepsilon_1, \varepsilon_2, \dots$ are independent Rademacher random variables, i.e., $\mathbb{P}(\varepsilon_t = 1) = \mathbb{P}(\varepsilon_t = -1) = \frac{1}{2}$, and $L > 0$ is the Lipschitz constant for all the losses.

Now, fix any algorithm for OCO. Clearly, its regret satisfies

$$\max_{\varepsilon_1, \dots, \varepsilon_T} \max_{\mathbf{u} \in \{\mathbf{v}_1, \mathbf{v}_2\}} R_T(\mathbf{u}) \geq \mathbb{E} \left[\max_{\mathbf{u} \in \{\mathbf{v}_1, \mathbf{v}_2\}} R_T(\mathbf{u}) \right]$$

where the expectation is with respect to the random draw of $\varepsilon_1, \dots, \varepsilon_T$. Moreover, since $\mathbb{E}[L_t(\mathbf{w})] = 0$ for all \mathbf{w} , we have

$$\mathbb{E} \left[\max_{\mathbf{u} \in \{\mathbf{v}_1, \mathbf{v}_2\}} R_T(\mathbf{u}) \right] = \mathbb{E} \left[\max_{\mathbf{u} \in \{\mathbf{v}_1, \mathbf{v}_2\}} \sum_{t=1}^T L_t(\mathbf{u}) \right] \quad (2)$$

Now, using the elementary identity $\max\{a, b\} = \frac{1}{2}(a + b + |a - b|)$ and Khintchine inequality (see, e.g., Cesa-Bianchi & Lugosi 2006, Lemma 8.2) we obtain that the right-hand side of (2) is equal to

$$\frac{L}{2} \mathbb{E} \left[\left\| \sum_{t=1}^T \varepsilon_t \mathbf{z}_0^\top (\mathbf{v}_1 - \mathbf{v}_2) \right\| \right] = \frac{LD}{2} \mathbb{E} \left[\left\| \sum_{t=1}^T \varepsilon_t \right\| \right] \geq \frac{LD}{4} \sqrt{2T} \quad (3)$$

where the equality is proven using $\mathbf{z}_0^\top (\mathbf{v}_1 - \mathbf{v}_2) = D$ due to our choice of \mathbf{z}_0 . This shows that we cannot expect the regret to grow slower than $LD\sqrt{T}$ —where D is the Euclidean diameter of \mathbb{X} and L is the Lipschitz constant of the loss—unless the two main parameters of our setting, that is the model space \mathbb{X} and the loss process ℓ_1, ℓ_2, \dots , enjoy some additional properties.

Lower bounds arguments based on the Khintchine inequality are rather common in online learning, see Luo et al. (2016, Theorem 1) for an example close to the one presented here. The effectiveness of stochastic loss sequences to prove tight lower bounds in OCO settings is not accidental. To gain a better understanding of the connections between stochastic and online learning, see Rakhlin & Sridharan (2014) who study online learning as a minimax problem.

The game-theoretic roots of OCO are described by Cesa-Bianchi & Lugosi (2006). Since then, the interface between sequential optimization, game theory, and statistics has been intensively explored in many works, see the tutorial and surveys by Shalev-Shwartz (2007), Hazan (2016), McMahan (2017), Orabona (2019).

2. ONLINE MIRROR DESCENT

We now introduce the most popular algorithm for OCO, Online Mirror Descent (or OMD for short). OMD is the online version of the Mirror Descent algorithm of Nemirovsky & Yudin (1983) for convex optimization. Mirror Descent is based on a generalization of projected gradient descent in which distances in the model space \mathbb{X} are not necessarily measured using the Euclidean norm. This allows to take advantage of specific geometrical properties that \mathbb{X} may have. To see how this is done, we start from the iterates $\mathbf{w}_{t+1} = \Pi_{\mathbb{X}}(\mathbf{w}_t - \eta_t \nabla F(\mathbf{w}_t))$ of projected gradient descent on a convex and differentiable objective $F: \mathbb{R}^d \rightarrow \mathbb{R}$, where

$\Pi_{\mathbb{X}}$ denotes the Euclidean projection onto \mathbb{X} and $\eta_t > 0$ is the *step size* at time t . The expression defining the iterates can be rewritten in an equivalent optimization form,

$$\mathbf{w}_{t+1} = \operatorname{argmin}_{\mathbf{w} \in \mathbb{X}} \frac{1}{2\eta_t} \|\mathbf{w} - \mathbf{w}_t\|_2^2 + \mathbf{w}^\top \nabla F(\mathbf{w}_t)$$

Mirror Descent replaces the Euclidean norm in the above equation with a generalized distance, or *divergence* D ,

$$\mathbf{w}_{t+1} = \operatorname{argmin}_{\mathbf{w} \in \mathbb{X}} \frac{1}{2\eta_t} D(\mathbf{w}, \mathbf{w}_t) + \mathbf{w}^\top \nabla F(\mathbf{w}_t) \quad (4)$$

Following Beck & Teboulle (2003), the divergences used by Mirror Descent are parameterized by *mirror map* functions $\psi : \mathbb{X} \rightarrow \mathbb{R}$ that are strictly convex and continuously differentiable on the interior of \mathbb{X} . Given such a ψ , the *Bregman divergence* $D_\psi : \mathbb{X} \times \operatorname{int} \mathbb{X} \rightarrow \mathbb{R}$ is defined by

$$D_\psi(\mathbf{u}, \mathbf{w}) = \psi(\mathbf{u}) - \psi(\mathbf{w}) - \nabla \psi(\mathbf{w})^\top (\mathbf{u} - \mathbf{w})$$

Note that D_ψ is not necessarily symmetric and, since ψ is strictly convex, is always nonnegative and equals zero only when $\mathbf{w} = \mathbf{u}$. When ψ is also twice differentiable, then Taylor's theorem shows that

$$D_\psi(\mathbf{u}, \mathbf{w}) = \frac{1}{2} (\mathbf{u} - \mathbf{w})^\top \nabla^2 \psi(\mathbf{z}) (\mathbf{u} - \mathbf{w})$$

for some \mathbf{z} on the line segment joining \mathbf{u} and \mathbf{w} . In other words, for ψ that are twice differentiable the divergence locally behaves like a squared Mahalanobis distance. Just like in the Euclidean case, we can also write the Mirror Descent update (4) using a Bregman projection (see, e.g., Bubeck 2015, Section 4).

The online version of Mirror Descent is now straightforward to obtain. In order to avoid considering iterates \mathbf{w}_t on the boundary of \mathbb{X} , where $D_\psi(\cdot, \mathbf{w}_t)$ is not defined, we restrict the argmin in (4) to a convex and nonempty subset $\mathbb{V} \subseteq \operatorname{int} \mathbb{X}$ and, consequently, measure the regret $R_T(\mathbf{u})$ only for $\mathbf{u} \in \mathbb{V}$. Let $\mathbf{w}_1 \in \mathbb{V}$ and fix a sequence $\eta_1 \geq \eta_2 \geq \dots > 0$ of step sizes. Now, for any sequence ℓ_1, ℓ_2, \dots of differentiable loss functions, the iterates of OMD are defined by

$$\mathbf{w}_{t+1} = \operatorname{argmin}_{\mathbf{w} \in \mathbb{V}} \frac{1}{2\eta_t} D_\psi(\mathbf{w}, \mathbf{w}_t) + \mathbf{w}^\top \mathbf{g}_t \quad (5)$$

where $\mathbf{g}_t = \nabla \ell_t(\mathbf{w}_t)$. The differentiability assumption for the losses ℓ_t can be relaxed to subdifferentiability, in which case \mathbf{g}_t is any element of the subdifferential of ℓ_t at \mathbf{w}_t . This is useful because some popular loss functions, like the hinge loss $\ell_t(\mathbf{w}_t) = \max\{0, 1 - y_t \mathbf{w}_t^\top \mathbf{x}_t\}$ for binary classification ($y_t \in \{-1, 1\}$), are not everywhere differentiable. In the rest of this survey, we use the same notation \mathbf{g}_t to denote $\nabla \ell_t(\mathbf{w}_t)$ or any subgradient of ℓ_t at \mathbf{w}_t , according to whether ℓ_t is differentiable or only subdifferentiable.

One might wonder why using a linear approximation $\mathbf{w}^\top \mathbf{g}_t$ instead of the loss $\ell_t(\mathbf{w})$ in the update of OMD. Indeed, the variant where $\ell_t(\mathbf{w})$ replaces $\mathbf{w}^\top \mathbf{g}_t$ is called *proximal point method* in the convex optimization literature and *implicit update* in the online learning literature, see Kivinen & Warmuth (1997), Kulis & Bartlett (2010) and also McMahan (2017, Section 6). Connections between implicit updates and optimistic updates in saddle point optimization problems were investigated by Mokhtari et al. (2019), see also Section 3 of this survey. We now look at two important choices for the mirror map.

Euclidean. If $\psi = \frac{1}{2} \|\cdot\|_2^2$, then $D_\psi(\mathbf{u}, \mathbf{w}) = \frac{1}{2} \|\mathbf{u} - \mathbf{w}\|_2^2$ and the OMD update become the online version of gradient descent (OGD) with Euclidean projection, $\mathbf{w}_{t+1} = \Pi_{\mathbb{V}}(\mathbf{w}_t - \eta_t \mathbf{g}_t)$.

Entropic. If \mathbb{X} is the simplex of probability distributions over $\{1, \dots, d\}$ and ψ is the negative entropy $\psi(\mathbf{w}) = \sum_i w_i \ln w_i$, then $D_\psi(\mathbf{u}, \mathbf{w}) = \sum_i u_i \ln \frac{u_i}{w_i}$ is the Kullback-Leibler divergence (or cross-entropy) and the OMD update becomes the Exponentiated Gradient (EG) algorithm of Kivinen & Warmuth (1997),

$$w_{t+1,i} = \frac{w_{t,i} e^{-\eta_t g_{t,i}}}{\sum_{j=1}^d w_{t,j} e^{-\eta_t g_{t,j}}} \quad i = 1, \dots, d$$

where $g_{t,i}$ are the components of the gradient \mathbf{g}_t . As in this case

$$\lim_{\mathbf{w} \rightarrow \text{bd } \mathbb{X}} \|\nabla \psi(\mathbf{w})\|_2 = \infty \quad (6)$$

where $\text{bd } \mathbb{X}$ is the boundary of \mathbb{X} , we can measure the regret $R_T(\mathbf{u})$ against $\mathbf{u} \in \mathbb{X}$ instead of restricting to $\mathbf{u} \in \mathbb{V} \subseteq \text{int } \mathbb{X}$. In fact, under the latter condition the update rule of OMD will never return a point on the boundary of \mathbb{X} . For instance, when losses are linear, $\sum_t \ell_t$ is always minimized at a \mathbf{u} located on a corner of the simplex, and we can show that EG (OMD with entropic mirror map) has vanishing regret with respect to any \mathbf{u} in the simplex \mathbb{X} , including the corners.

We now look at a different interpretation of the update rule of OMD. First, let a differentiable function $f : \mathbb{X} \rightarrow \mathbb{R}$ be μ -strongly convex on $\mathbb{V} \subseteq \text{int } \mathbb{X}$ with respect to a norm $\|\cdot\|$ if for all $\mathbf{u}, \mathbf{v} \in \mathbb{V}$ we have that $f(\mathbf{u}) \geq f(\mathbf{v}) + \nabla f(\mathbf{v})^\top (\mathbf{u} - \mathbf{v}) + \frac{\mu}{2} \|\mathbf{u} - \mathbf{v}\|^2$. Mirror maps ψ that are μ -strongly convex with respect to a norm $\|\cdot\|$ define Bregman divergences that grow faster than the square of the same norm, i.e., $D_\psi(\mathbf{u}, \mathbf{w}) \geq \frac{\mu}{2} \|\mathbf{u} - \mathbf{w}\|^2$.

For strongly convex mirror maps, an equivalent way of writing (5) is $\mathbf{w}_{t+1} = \nabla \psi_V^*(\nabla \psi_V(\mathbf{w}_t) - \eta_t \mathbf{g}_t)$, where the function ψ_V is the restriction of ψ to \mathbb{V} , i.e., $\psi_V(\mathbf{w}) = \psi(\mathbf{w})$ if $\mathbf{w} \in \mathbb{V}$ and $\psi_V(\mathbf{w}) = \infty$ otherwise. The function $\psi_V^* : \mathbb{R}^d \rightarrow \mathbb{R}$ is the *Fenchel conjugate* of ψ_V , defined by

$$\psi_V^*(\boldsymbol{\theta}) = \max_{\mathbf{w} \in \mathbb{R}^d} \mathbf{w}^\top \boldsymbol{\theta} - \psi_V(\mathbf{w}) \quad (7)$$

Strong convexity ensures that ψ_V^* is differentiable and that $\nabla \psi_V^*$ is the functional inverse of $\nabla \psi_V$. The function ψ_V maps the iterates \mathbf{w}_t to the dual space of gradients where a gradient step is performed. The inverse function ψ_V^* maps back to the primal space of iterates. See Figure 1 for a picture illustrating the OMD update. Going back to the previous examples, we see that $\psi = \frac{1}{2} \|\cdot\|_2^2 = \psi^*$. When ψ is the negative entropy we instead have that $\psi^*(\boldsymbol{\theta}) = \ln \left(\sum_{i=1}^d e^{\theta_i} \right)$.

We now move on to analyze the regret of OMD, and show us how the change of geometry caused by the choice of the strongly convex mirror map ψ affects the algorithm's performance. The first step in bounding the regret $R_T(\mathbf{u})$ consists in linearizing the loss using convexity, $\ell_t(\mathbf{w}_t) - \ell_t(\mathbf{u}) \leq \mathbf{g}_t^\top (\mathbf{w}_t - \mathbf{u})$, recall that $\mathbf{g}_t = \nabla \ell_t(\mathbf{w}_t)$ for differentiable losses. We call *linearized regret* the resulting upper bound on $R_T(\mathbf{u})$. Next, using the properties of the optimization form for the OMD iterates (4), and the μ -strong convexity of the mirror map with respect to $\|\cdot\|$, it takes a few steps to prove that

$$\eta_t \mathbf{g}_t^\top (\mathbf{w}_t - \mathbf{u}) \leq D_\psi(\mathbf{u}, \mathbf{w}_t) - D_\psi(\mathbf{u}, \mathbf{w}_{t+1}) + \frac{\eta_t^2}{2\mu} \|\mathbf{g}_t\|_*^2$$

where $\|\cdot\|_*$ is the dual norm of $\|\cdot\|$. Now, dividing both sides by $\eta_t > 0$ and summing over




Figure 1

A step of Online Mirror Descent (OMD) run with ψ_V (the restriction to the model space $V \subseteq \mathbb{X}$ of a strongly convex mirror map ψ). The function ψ_V^* is the Fenchel conjugate of ψ_V . The vector \mathbf{g}_t denotes the loss gradient $\nabla \ell_t(\mathbf{w}_t)$ (or any subgradient of ℓ_t at \mathbf{w}_t) and $\eta_t > 0$ is a variable stepsize.

$t = 1, \dots, T$ gives the following chain of inequalities

$$\begin{aligned}
R_T(\mathbf{u}) &= \sum_{t=1}^T (\ell_t(\mathbf{w}_t) - \ell_t(\mathbf{u})) \\
&\leq \sum_{t=1}^T \mathbf{g}_t^\top (\mathbf{w}_t - \mathbf{u}) \quad (\text{linearized regret}) \\
&\leq \sum_{t=1}^T \left(\frac{D_\psi(\mathbf{u}, \mathbf{w}_t)}{\eta_t} - \frac{D_\psi(\mathbf{u}, \mathbf{w}_{t+1})}{\eta_t} \right) + \frac{1}{2\mu} \sum_{t=1}^T \eta_t \|\mathbf{g}_t\|_*^2 \\
&= \frac{D_\psi(\mathbf{u}, \mathbf{w}_1)}{\eta_1} - \frac{D_\psi(\mathbf{u}, \mathbf{w}_{T+1})}{\eta_{T+1}} + \sum_{t=1}^{T-1} \left(\frac{1}{\eta_{t+1}} - \frac{1}{\eta_t} \right) D_\psi(\mathbf{u}, \mathbf{w}_{t+1}) + \frac{1}{2\mu} \sum_{t=1}^T \eta_t \|\mathbf{g}_t\|_*^2 \\
&\leq \frac{D^2}{\eta_1} + \left(\frac{1}{\eta_T} - \frac{1}{\eta_1} \right) D^2 + \frac{1}{2\mu} \sum_{t=1}^T \eta_t \|\mathbf{g}_t\|_*^2 \quad (\text{where } D^2 = \max_{\mathbf{u}, \mathbf{w} \in V} D_\psi(\mathbf{u}, \mathbf{w})) \\
&= \frac{D^2}{\eta_T} + \frac{1}{2\mu} \sum_{t=1}^T \eta_t \|\mathbf{g}_t\|_*^2 \tag{8}
\end{aligned}$$

We can now set $\eta_t = D \sqrt{\frac{\mu}{\sum_{s=1}^t \|\mathbf{g}_s\|_*^2}}$ and obtain

$$R_T(\mathbf{u}) \leq 2D \sqrt{\frac{1}{\mu} \sum_{t=1}^T \|\mathbf{g}_t\|_*^2} \tag{9}$$

Equipped with this result, we can see how the choice of the mirror map affects OMD performance. Consider first OGD, where V is the closed Euclidean ball of diameter U and $\psi = \frac{1}{2} \|\cdot\|_2^2$. If $\|\mathbf{g}_t\|_\infty \leq G$ for some $G > 0$, then $\|\mathbf{g}_t\|_*^2 = \|\mathbf{g}_t\|_2^2 \leq G^2 d$ and OGD has a regret bound of the form $R_T(\mathbf{u}) \leq 2UG\sqrt{dT}$. Under these assumptions, losses have Lipschitz constant $L = \max_t \|\mathbf{g}_t\|_2 \leq G\sqrt{d}$, and so this regret bound matches the lower bound (3) of Subsection 1.3 up to constant factors. Therefore, in the Euclidean domain, OGD has essentially optimal dependence on time, diameter of the model space, and Lipschitz constant of the losses.

Next, consider EG, where \mathbb{V} is the probability simplex and $\psi(\mathbf{w}) = \sum_i w_i \ln w_i$, which is 1-strongly convex with respect to $\|\cdot\|_1$. Here, we run into a problem because the diameter of the simplex is unbounded when measured using the cross-entropy (i.e., the Bregman divergence D_ψ corresponding to the entropic mirror map). This prevents us from obtaining a constant upper bound D^2 on $\max_t D_\psi(\mathbf{u}, \mathbf{w}_t)$. We can fix this by choosing a constant step size η , which allows us to transform the regret guarantee (8) into the following alternative bound,

$$R_T(\mathbf{u}) \leq \frac{D_\psi(\mathbf{u}, \mathbf{w}_1)}{\eta} + \frac{\eta}{2\mu} \sum_{t=1}^T \|\mathbf{g}_t\|_*^2 \quad (10)$$

Under the same assumption $\|\mathbf{g}_t\|_\infty \leq G$ as before, we get $\|\mathbf{g}_t\|_*^2 = \|\mathbf{g}_t\|_\infty^2 \leq G^2$. Hence, choosing $\mathbf{w}_1 = (1/d, \dots, 1/d)$ as first iterate so that $D_\psi(\mathbf{u}, \mathbf{w}_1) \leq \ln d$, and setting $\eta = \sqrt{(2 \ln d)/(G^2 T)}$, gives $R_T(\mathbf{u}) \leq G\sqrt{(T \ln d)/2}$. Using the squared Euclidean norm in the same setting would instead give $R_T(\mathbf{u}) \leq 2G\sqrt{dT}$. Note that for linear losses the quantity $G\sqrt{(T \ln d)/2}$ is asymptotically minimax, *including constants* (see Cesa-Bianchi & Lugosi 2006, Corollary 8.3). This shows the importance of matching the mirror map to the geometry of the model space. From a practical point of view, the logarithmic dependence on d in the regret guarantees that EG is robust to a large number of irrelevant features.

When the convex losses ℓ_t are induced by pairs (\mathbf{x}_t, y_t) , as in $\ell_t(\mathbf{w}) = (\mathbf{w}^\top \mathbf{x}_t - y_t)^2$, then loss gradients $\nabla \ell_t(\mathbf{w})$ are proportional to data points \mathbf{x}_t . In this case, bounds on $R_T(\mathbf{u})$ for OGD and EG depend on different products of dual norms: $(\max_t \|\mathbf{x}_t\|_2) \|\mathbf{u}\|_2$ for OGD and $(\max_t \|\mathbf{x}_t\|_\infty) \|\mathbf{u}\|_1$ for EG. Since $\|\cdot\|_\infty \leq \|\cdot\|_2 \leq \|\cdot\|_1$, neither bound dominates. However, for sparse \mathbf{u} , $\|\mathbf{u}\|_1$ approaches $\|\mathbf{u}\|_\infty$ and EG performs better than OGD.

The analysis of OGD can be easily adapted to derive a regret bound for the Perceptron algorithm of Rosenblatt (1958) for binary classification, without need of assuming linear separability of the data sequence (see Freund & Schapire 1999). In the special case of separable data sequences, the regret bound reduces to the result originally proven in the well-known Perceptron Convergence Theorem (Block 1962, Novikoff 1963).

If the loss functions are known to be μ -strongly convex and L -Lipschitz in \mathbb{V} , then the step size of OGD can be set more aggressively to $\eta_t = \frac{1}{\mu t}$ so to exploit the curvature. The resulting bound on the regret is

$$R_T(\mathbf{u}) \leq \frac{L^2}{2\mu} \ln(T+1) \quad \forall \mathbf{u} \in \mathbb{V}$$

Note that there is no dependence on the diameter of the model space in this bound, but only on the gradients of the losses (through the Lipschitz constant L). Many machine learning algorithms, including Support Vector Machines (Cortes & Vapnik 1995), can be trained using stochastic optimization over a training set of data points $\mathbf{z}_1, \dots, \mathbf{z}_m$ to minimize a strongly convex functional F , see (1). A prime example of a stochastic gradient descent algorithm is OGD run on the sequence of strongly convex losses $\ell_t = F(\cdot, \mathbf{Z}_t)$, where \mathbf{Z}_t is drawn at random from the dataset. The regret analysis of OGD with strongly convex losses can be used to obtain rates of convergence to the minimum of F in (1). This is done, for example, to analyze the Pegasos algorithm of Shalev-Shwartz et al. (2011).

The EG algorithm is an instance of the multiplicative update method, a technique that found applications in computer science (Littlestone & Warmuth 1994), game theory (Fudenberg & Levine 1995), information theory, statistics, and other disciplines (see, e.g., Cesa-Bianchi & Lugosi 2006, Arora et al. 2012b). In the special case of linear losses with

uniformly bounded coefficients, EG is known as the Hedge algorithm (Freund & Schapire 1997) and the corresponding setting is known as prediction with expert advice (Cesa-Bianchi et al. 1997).

2.1. The AdaGrad algorithm

A variant of OMD that has become of widespread use as a stochastic gradient descent algorithm for training neural networks is the AdaGrad algorithm, independently introduced by McMahan & Streeter (2010) and Duchi et al. (2011). For simplicity, we look at the so-called diagonal version of AdaGrad, which uses a coordinate-dependent step size.

Let \mathbb{V} be the hyperrectangle $[a_1, b_1] \times \dots \times [a_d, b_d]$ and $D_i = b_i - a_i$ for $i = 1, \dots, d$. AdaGrad runs OMD with mirror map $\psi = \frac{1}{2} \|\cdot\|^2$ and projection onto \mathbb{V} . The iterates, including the projection step, can be written as $w_{t+1,i} = \max\{\min\{w_{t,i} - \eta_{t,i} g_{t,i}, a_i\}, b_i\}$ for $i = 1, \dots, d$ and an arbitrary initial point $\mathbf{w}_1 \in \mathbb{V}$. The components of the step size are chosen as

$$\eta_{t,i} = \frac{D_i}{\sqrt{2 \sum_{s=1}^t g_{s,i}^2}}$$

The regret analysis is straightforward: after linearizing the losses so that $\ell_t(\mathbf{w}_t) - \ell_t(\mathbf{u}) \leq \mathbf{g}_t^\top (\mathbf{w}_t - \mathbf{u})$, one can simply perform the standard OMD analysis independently for each coordinate. By applying the bound (9) on each coordinate, and then summing over coordinates, we obtain

$$R_T(\mathbf{u}) \leq \sum_{i=1}^d D_i \sqrt{2 \sum_{t=1}^T g_{t,i}^2} \quad (11)$$

We can compare this bound to (9) with $\mu = 1$ and $\|\cdot\|_* = \|\cdot\|_2$. For simplicity, we take \mathbb{V} to be the hypercube with $D_i = 1$, so that D in (9) is equal to \sqrt{d} . Using Cauchy-Schwartz,

$$\sqrt{\sum_{t=1}^T \|\mathbf{g}_t\|_2^2} \leq \sum_{i=1}^d \sqrt{\sum_{t=1}^T g_{t,i}^2} \leq \sqrt{d} \sqrt{\sum_{t=1}^T \|\mathbf{g}_t\|_2^2}$$

where (9) appears in the right-hand side and (11) appears in the middle. Hence, ignoring the fact that the sequence of realized subgradients \mathbf{g}_t is different for the two algorithms, with this choice of \mathbb{V} AdaGrad can gain up to a factor of \sqrt{d} in the regret bound with respect to plain OMD. Compared to the EG algorithm, which uses an entropic mirror map on the probability simplex, here the advantage is brought by a coordinate-dependent step size, which exploits the decomposability of the regret across the d coordinates granted by the geometry of \mathbb{V} .

Note that the specific choice of the step size makes the algorithm independent with respect to rescalings of the coordinates (Orabona & Pál 2018). This property is especially useful in neural network training, where the range of gradient components may vary a lot across the different layers.

3. FOLLOW THE REGULARIZED LEADER

A very natural online learning strategy for the OCO setting is Follow the Leader (FTL), which corresponds to predicting with the model minimizing the sum of the losses observed

so far,

$$\mathbf{w}_{t+1} = \operatorname{argmin}_{\mathbf{w} \in \mathbb{V}} \sum_{s=1}^t \ell_s(\mathbf{w})$$

Here $\mathbf{w}_1 \in \mathbb{V}$ is any convex, closed, and nonempty subset of \mathbb{R}^d . When losses are strongly convex and Lipschitz in \mathbb{V} , FTL achieves $R_T(\mathbf{u}) = \mathcal{O}(\ln T)$ for all $\mathbf{u} \in \mathbb{V}$ (McMahan 2017, Section 3.7). However, the curvature of each loss function is necessary to obtain a nontrivial performance. When losses are linear, FTL provably incurs a regret that grows linearly in T (see, e.g., Shalev-Shwartz 2012, Example 2.2). This is caused by the intrinsic instability of the algorithm: in a nonstochastic setting, one can design the loss sequence so that the trajectory $\mathbf{w}_1, \mathbf{w}_2, \dots$ of FTL models oscillates wildly in \mathbb{V} , a behavior that the adversary can exploit to increase the regret. Similarly to OMD, which achieves stability by forcing \mathbf{w}_{t+1} to be not too far away from \mathbf{w}_t , see (5), FTL can be stabilized by adding a strictly convex regularization function. The resulting algorithm, appropriately called Follow the Regularized Leader (FTRL), is a close relative of OMD. Indeed, the regularization functions used by FTRL are formally equivalent to OMD's mirror maps ψ . As FTRL is not formulated as a gradient descent method, we replace OMD's step sizes η_t with time-dependent regularizers ψ_t . For any sequence ψ_1, ψ_2, \dots of strictly convex regularizers, FTRL iterates are defined by

$$\mathbf{w}_{t+1} = \operatorname{argmin}_{\mathbf{w} \in \mathbb{V}} \psi_{t+1}(\mathbf{w}) + \sum_{s=1}^t \ell_s(\mathbf{w}) \quad (12)$$

where ℓ_1, ℓ_2, \dots is an arbitrary sequence of convex losses. From the viewpoint of regret minimization, we know that $\ell_s(\mathbf{w})$ can be replaced by $\mathbf{g}_s^\top \mathbf{w}$ (recall the linearization step in Section 2), where \mathbf{g}_s is the gradient (or subgradient) of $\ell_s(\mathbf{w}_s)$. If we do that in (12), we obtain

$$\mathbf{w}_{t+1} = \operatorname{argmin}_{\mathbf{w} \in \mathbb{V}} \psi_{t+1}(\mathbf{w}) + \sum_{s=1}^t \mathbf{g}_s^\top \mathbf{w} \quad (13)$$

This is the version of FTRL we study in the rest of this section. Let $\psi_{\mathbb{V},t}$ be the restriction of ψ_t to \mathbb{V} (see Section 2). If the regularizers $\psi_{\mathbb{V},t}$ are all strongly convex, then \mathbf{w}_{t+1} in (13) has the closed form

$$\mathbf{w}_{t+1} = \nabla \psi_{\mathbb{V},t+1}^* \left(- \sum_{s=1}^t \mathbf{g}_s \right)$$

where $\psi_{\mathbb{V},t+1}^*$ is the Fenchel conjugate of $\psi_{\mathbb{V},t+1}$ (differentiability of $\psi_{\mathbb{V},t+1}^*$ is guaranteed by the strong convexity of $\psi_{\mathbb{V},t+1}$). By letting $\boldsymbol{\theta}_{t+1} = -(\mathbf{g}_1 + \dots + \mathbf{g}_t)$, the FTRL iterates can be written as $\mathbf{w}_{t+1} = \nabla \psi_{\mathbb{V},t+1}^*(\boldsymbol{\theta}_{t+1})$, see Figure 2. By comparison, the OMD iterates for strongly convex mirror maps are written as $\mathbf{w}_{t+1} = \nabla \psi_{\mathbb{V}}^*(\boldsymbol{\theta}'_{t+1})$ where $\boldsymbol{\theta}'_{t+1} = \nabla \psi_{\mathbb{V}}(\mathbf{w}_t) - \eta_t \mathbf{g}_t$.

Note that FTRL keeps a state variable $\boldsymbol{\theta}_t$ in the dual space of gradients. This is mapped to the primal space of iterates every time a prediction is needed. On the other hand, OMD keeps its state \mathbf{w}_t in the primal space of iterates. This is then mapped to the dual space of gradients every time an update must be computed. Another difference is that gradients are all equally weighted in FTRL, whereas in OMD each gradient \mathbf{g}_t is weighted by a potentially different step size η_t .

FTRL was originally introduced by Abernethy et al. (2008), although the key ideas are already contained in (Shalev-Shwartz 2007). The version with linearized losses and $\psi_t = \psi = \frac{\eta}{2} \|\cdot\|_2^2$ was introduced by Zinkevich (2004) as a “lazy update” variant of OMD.




Figure 2

A step of Follow the Regularized Leader (FTRL) run with linearized losses and time-dependent regularizers $\psi_{V,t}$ (these are the restrictions to the model space $V \subseteq \mathbb{X}$ of strongly convex regularizers ψ_t). The functions $\psi_{V,t}^*$ are the Fenchel conjugate of $\psi_{V,t}$. The vectors \mathbf{g}_t denote the loss gradients $\nabla \ell_t(\mathbf{w}_t)$ (or any subgradient of ℓ_t at \mathbf{w}_t), while the state variables $\boldsymbol{\theta}_t$ are simply the sum of all past loss gradients, $\boldsymbol{\theta}_t = -(\mathbf{g}_1 + \dots + \mathbf{g}_{t-1})$.

In order to understand whether the two algorithms can produce the same sequence of iterates, we focus on the case when $\eta_t = \eta$ in OMD and $\psi_t = \psi/\eta$ in FTRL, for all t . Then we define the iterates of both OMD and FTRL through a common optimization problem $\mathbf{w}_{t+1} = \operatorname{argmin}_{\mathbf{w} \in V} F_{t+1}(\mathbf{w})$, where

$$F_{t+1}(\mathbf{w}) = D_\psi(\mathbf{w}, \mathbf{w}_t) + \eta \mathbf{g}_t^\top \mathbf{w} \quad (\text{OMD})$$

$$F_{t+1}(\mathbf{w}) = \psi(\mathbf{w}) + \sum_{s=1}^t \mathbf{g}_s^\top \mathbf{w} \quad (\text{FTRL})$$

When V is such that \mathbf{w}_{t+1} satisfies $\nabla F_{t+1}(\mathbf{w}_{t+1}) = \mathbf{0}$ for both instances of F_{t+1} , a quick computation shows us that for both choices of F_{t+1} the gradient vanishes when $\nabla \psi(\mathbf{w}_{t+1}) = -\eta \boldsymbol{\theta}_{t+1}$. Due to the strong convexity of ψ , this is equivalent to $\mathbf{w}_{t+1} = \nabla \psi^*(-\eta \boldsymbol{\theta}_{t+1})$. Under these conditions, the two algorithms produce the same sequence of iterates. A concrete setting in which the two algorithms become identical is $V \equiv \mathbb{R}^d$ and $\psi = \frac{1}{2} \|\cdot\|_2^2$.

FTRL enjoys a regret bound similar to the one we stated for OMD. Let $\psi : V \rightarrow \mathbb{R}$ be a μ -strongly convex function with respect to a norm $\|\cdot\|$. Since FTRL is invariant to positive constants added to the regularizers, without loss of generality we may assume that $\min_{\mathbf{u} \in V} \psi(\mathbf{u}) = 0$. For scaling factors $\beta_1 \geq \beta_2 \geq \dots > 0$, let $\psi_t = \psi/\beta_t$. Then, for any sequence ℓ_1, ℓ_2, \dots of convex losses, the regret of FTRL satisfies

$$R_T(\mathbf{u}) \leq \frac{\psi(\mathbf{u})}{\beta_T} + \frac{1}{2\mu} \sum_{t=1}^T \beta_t \|\mathbf{g}_t\|_*^2 \quad \forall \mathbf{u} \in V \quad (14)$$

This is very similar to the OMD bound (8). However, unlike OMD where the step size η_t is used in the update $\mathbf{w}_t \rightarrow \mathbf{w}_{t+1}$, in the above formulation FTRL uses the scaling factor β_t in the update $\mathbf{w}_{t-1} \rightarrow \mathbf{w}_t$. In particular, in FTRL β_t cannot depend on \mathbf{g}_t .

Given the similarities between OMD and FTRL, are there settings in which the latter should be preferred over the former? Consider for instance the analysis of the EG algorithm (i.e., OMD run with entropic mirror map and V set to the probability simplex). In this case we can not use OMD with a variable step size because of the unboundedness of the Bregman

divergence. When FTRL is applied to the same setting, the regret bound specializes to

$$R_T(\mathbf{u}) \leq \frac{\ln d}{\beta_T} + \frac{G^2}{2} \sum_{t=1}^T \beta_t$$

where, we recall it, G upper bounds $\|\mathbf{g}_t\|_\infty$ for all t . Taking $\beta_t = \sqrt{(\ln d)/(G^2 t)}$ then gives $R_T(\mathbf{u}) \leq 2G\sqrt{T \ln d}$. Up to constants, this is the same as the bound we obtain for OMD when the step size is tuned with prior knowledge of the number of rounds T .

3.1. Online Newton Step

As mentioned at the beginning of this section, strongly convex losses are an easy case for online learning. A simple algorithm like FTL achieves logarithmic regret on any sequence of such losses whenever the strong convexity coefficients in the loss sequence are bounded away from zero. A natural question is then whether logarithmic regret is possible for convex loss functions that are not strongly convex but also not linear.

For any symmetric and positive semidefinite matrix M , introduce the semi-norm $\|\cdot\|_M$ such that $\|\mathbf{w}\|_M^2 = \mathbf{w}^\top M \mathbf{w}$. As it turns out, the right curvature property sufficient to guarantee logarithmic regret is the following:

$$\ell_t(\mathbf{u}) \geq \ell_t(\mathbf{w}) + \mathbf{g}^\top (\mathbf{u} - \mathbf{w}) + \frac{\lambda}{2} \|\mathbf{u} - \mathbf{w}\|_{\mathbf{g}\mathbf{g}^\top}^2 \quad \mathbf{u}, \mathbf{w} \in \mathbb{V} \quad (15)$$

for some $\lambda > 0$, where $\mathbf{g} = \nabla \ell_t(\mathbf{w})$ (or any subgradient if ℓ_t is only subdifferentiable). In words, we require ℓ_t to be strongly convex only in the direction of its gradient (or in the direction of some of its subgradients). For example, the square loss $\ell_t(\mathbf{w}) = \frac{1}{2} (\mathbf{w}^\top \mathbf{x}_t - y_t)^2$ satisfies property (15) for $\lambda \leq \frac{1}{8C^2}$ whenever $|\mathbf{w}^\top \mathbf{x}_t|, |y_t| \leq C$ (Hazan et al. 2007, Lemma 3). Also, the logistic loss $\ell_t(\mathbf{w}) = \ln(1 + \exp(-\mathbf{w}^\top \mathbf{x}_t))$ satisfies property (15) when \mathbf{w} belongs to an Euclidean ball of fixed radius. More in general, any loss ℓ_t such that $e^{-\alpha \ell_t}$ is concave in \mathbb{V} for some $\alpha > 0$ satisfies property (15) with $\lambda \leq \frac{1}{2} \min\{\frac{1}{GD}, \alpha\}$, where D is the Euclidean diameter of \mathbb{V} and G is a bound on $\max_{\mathbf{w} \in \mathbb{V}} \|\nabla \ell_t(\mathbf{w})\|$. Such losses are said to be α -exp-concave in \mathbb{V} (Kivinen & Warmuth 1999). If ℓ_t is twice differentiable, then α -exp-concavity in \mathbb{V} is equivalent to $\nabla^2 \ell_t(\mathbf{w}) - \alpha \nabla \ell_t(\mathbf{w}) \nabla \ell_t(\mathbf{w})^\top$ being positive semidefinite for all $\mathbf{w} \in \mathbb{V}$.

Hazan et al. (2007, Theorem 6) prove a $\mathcal{O}(d \ln T)$ regret bound on FTL for all sequences of loss functions ℓ_t satisfying (15) for some $\lambda > 0$. They do so by introducing the quadratic approximation of each loss function $\hat{\ell}_t(\mathbf{w}) = \ell_t(\mathbf{w}_t) + \mathbf{g}_t^\top (\mathbf{w} - \mathbf{w}_t) + \frac{\lambda}{2} \|\mathbf{w} - \mathbf{w}_t\|_{\mathbf{g}_t \mathbf{g}_t^\top}^2$, where $\mathbf{g}_t = \nabla \ell_t(\mathbf{w}_t)$ and $\mathbf{w}_1, \mathbf{w}_2, \dots \in \mathbb{V}$ are defined by

$$\mathbf{w}_{t+1} = \operatorname{argmin}_{\mathbf{w} \in \mathbb{V}} \sum_{s=1}^t \hat{\ell}_s(\mathbf{w}) \quad (16)$$

(\mathbf{w}_1 can be defined arbitrarily). Observe that (16) is just FTL run on the losses $\hat{\ell}_t$. Moreover, $\hat{\ell}_t(\mathbf{w}_t) = \ell_t(\mathbf{w}_t)$ and, because of (15), $\hat{\ell}_t(\mathbf{u}) \leq \ell_t(\mathbf{u})$ for all $\mathbf{u} \in \mathbb{V}$. This implies that the regret of FTL with respect to the original losses ℓ_t satisfies

$$R_T(\mathbf{u}) \leq \sum_{t=1}^T \hat{\ell}_t(\mathbf{w}_t) - \sum_{t=1}^T \hat{\ell}_t(\mathbf{u}) \quad \forall \mathbf{u} \in \mathbb{V}$$

The rest of the proof uses the special properties of the functions $\widehat{\ell}_t$ to derive the desired $\mathcal{O}(d \ln T)$ bound on the regret. Hazan et al. (2007, Lemma 4) also prove that (16) is equivalent to

$$\begin{aligned}\mathbf{w}'_{t+1} &= S_t^+ \sum_{s=1}^t \left(\mathbf{g}_s^\top \mathbf{w}_s - \frac{1}{\lambda} \right) \mathbf{g}_s \\ \mathbf{w}_{t+1} &= \operatorname{argmin}_{\mathbf{w} \in \mathbb{V}} \|\mathbf{w} - \mathbf{w}'_{t+1}\|_{S_t}\end{aligned}\tag{17}$$

where S_t^+ denotes the Moore-Penrose pseudoinverse of $S_t = \mathbf{g}_1 \mathbf{g}_1^\top + \dots + \mathbf{g}_t \mathbf{g}_t^\top$.

A similar $\mathcal{O}(d \ln T)$ bound can be proven through a more general proof, this time using the FTRL framework with linearized losses. Define the regularizers ψ_1, ψ_2, \dots given by

$$\psi_1(\mathbf{w}) = \frac{\|\mathbf{w}\|_2^2}{2} \quad \text{and} \quad \psi_{t+1}(\mathbf{w}) = \psi_t(\mathbf{w}) + \frac{\lambda}{2} \|\mathbf{w}_t - \mathbf{w}\|_{\mathbf{g}_t \mathbf{g}_t^\top}^2 \tag{18}$$

where $\mathbf{g}_t = \nabla \ell_t(\mathbf{w}_t)$ and the \mathbf{w}_t are the FTRL iterates defined in (13),

$$\mathbf{w}_{t+1} = \operatorname{argmin}_{\mathbf{w} \in \mathbb{V}} \psi_{t+1}(\mathbf{w}) + \sum_{s=1}^t \mathbf{g}_s^\top \mathbf{w} \tag{19}$$

An equivalent and more explicit form for the iterates (19), which brings out the similarity with (17), is

$$\begin{aligned}\mathbf{w}'_{t+1} &= A_t^{-1} \sum_{s=1}^t (\mathbf{g}_s^\top \mathbf{w}_s - 1) \mathbf{g}_s \\ \mathbf{w}_{t+1} &= \operatorname{argmin}_{\mathbf{w} \in \mathbb{V}} \|\mathbf{w} - \mathbf{w}'_{t+1}\|_{A_t}\end{aligned}\tag{20}$$

where $A_0 = I$ and $A_t = A_{t-1} + \lambda \mathbf{g}_t \mathbf{g}_t^\top$.

Each regularizer ψ_t of the form (18) is 1-strongly-convex with respect to the norm $\|\cdot\|_{A_{t-1}}$. Using this property, and the special recursive form of these regularizers, one can prove for this algorithm a regret bound of the following form: for any sequence ℓ_1, ℓ_2, \dots of losses satisfying (15),

$$R_T(\mathbf{u}) \leq \frac{\|\mathbf{u}\|_2^2}{2} + \frac{1}{2} \sum_{t=1}^T \|\mathbf{g}_t\|_{A_t^{-1}}^2 \leq \frac{\|\mathbf{u}\|_2^2}{2} + \frac{d}{2\lambda} \ln \left(1 + \frac{\lambda G T}{d} \right) \tag{21}$$

where $\max_t \|\mathbf{g}_t\|_2 \leq G$ and we used a standard majorization for bounding the sum of terms $\|\mathbf{g}_t\|_{A_t^{-1}}^2$ (see, e.g., Cesa-Bianchi & Lugosi 2006, Lemma 11.11 and Theorem 11.7).

One may also wonder whether the same logarithmic regret bound for losses satisfying (15) could be achieved using OMD instead of FTL/FTRL as we did in this section. Noting that, for M symmetric and positive definite, the divergence associated with $\psi = \frac{1}{2} \|\cdot\|_M^2$ is $D(\mathbf{u}, \mathbf{w}) = \frac{1}{2} \|\mathbf{u} - \mathbf{w}\|_M^2$, we may introduce the following instance of OMD

$$\mathbf{w}_{t+1} = \operatorname{argmin}_{\mathbf{w} \in \mathbb{V}} \frac{1}{2} \|\mathbf{w} - \mathbf{w}_t\|_{A_t}^2 + \mathbf{w}^\top \mathbf{g}_t$$

where $A_0 = I$, $A_t = A_{t-1} + \lambda \mathbf{g}_t \mathbf{g}_t^\top$, and $\mathbf{g}_t = \nabla \ell(\mathbf{w}_t)$. The closed-form expression is

$$\begin{aligned}\mathbf{w}'_{t+1} &= A_t^{-1} \mathbf{g}_t, \\ \mathbf{w}_{t+1} &= \operatorname{argmin}_{\mathbf{w} \in \mathbb{V}} \|\mathbf{w} - \mathbf{w}'_{t+1}\|_{A_t}\end{aligned}\tag{22}$$

Note that this algorithm is rather different from both (17) and (20). Its regret analysis is relatively simple, although it is not derived as a special case of the general OMD analysis (which uses constant mirror maps). Fix any sequence ℓ_1, ℓ_2, \dots of loss functions satisfying (15). Since \mathbf{w}_{t+1} is the projection of \mathbf{w}'_{t+1} onto \mathbb{V} , the update (22) ensures that

$$\|\mathbf{w}_{t+1} - \mathbf{u}\|_{A_t}^2 \leq \|\mathbf{w}'_{t+1} - \mathbf{u}\|_{A_t}^2 = \|\mathbf{w}_t - \mathbf{u}\|_{A_t}^2 + \mathbf{g}_t^\top A_t^{-1} \mathbf{g}_t - 2\mathbf{g}_t^\top (\mathbf{w}_t - \mathbf{u}) \quad (23)$$

for all $\mathbf{u} \in \mathbb{V}$. By the curvature property (15), we then have that

$$\begin{aligned} 2R_T(\mathbf{u}) &\leq 2 \sum_{t=1}^T \left(\mathbf{g}_t^\top (\mathbf{w}_t - \mathbf{u}) - \lambda \|\mathbf{w}_t - \mathbf{u}\|_{\mathbf{g}_t \mathbf{g}_t^\top}^2 \right) \\ &\leq \sum_{t=1}^T \left(\|\mathbf{g}_t\|_{A_t^{-1}} + \|\mathbf{w}_t - \mathbf{u}\|_{A_t}^2 - \|\mathbf{w}_{t+1} - \mathbf{u}\|_{A_t}^2 - \lambda \|\mathbf{w}_t - \mathbf{u}\|_{\mathbf{g}_t \mathbf{g}_t^\top}^2 \right) \quad (\text{using (23)}) \\ &= \sum_{t=1}^T \|\mathbf{g}_t\|_{A_t^{-1}}^2 + \|\mathbf{u}\|_{A_0}^2 + \sum_{t=1}^T (\mathbf{w}_t - \mathbf{u})^\top (A_t - A_{t-1} - \lambda \mathbf{g}_t \mathbf{g}_t^\top) (\mathbf{w}_t - \mathbf{u}) \\ &= \sum_{t=1}^T \|\mathbf{g}_t\|_{A_t^{-1}}^2 + \|\mathbf{u}\|_2^2 \quad (\text{using the definition of } A_t) \end{aligned}$$

Now note that the the above is exactly equivalent to the bound (21) which we proved for FTRL.

3.2. Online linear regression

An important special case of online learning is (unconstrained) online linear regression, where $\mathbb{V} = \mathbb{R}^d$ and the losses $\ell_t(\mathbf{w}) = \frac{1}{2} (\mathbf{w}^\top \mathbf{x}_t - y_t)^2$ are induced by an arbitrary and deterministic sequence $(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots$ of data points $\mathbf{x}_t \in \mathbb{R}^d$ and values $y_t \in \mathbb{R}$. The online version of the classical Ridge Regression algorithm by Hoerl & Kennard (2000) is an instance of FTRL without linearized losses, see (12),

$$\mathbf{w}_{t+1} = \underset{\mathbf{w} \in \mathbb{R}^d}{\operatorname{argmin}} \frac{1}{2} \|\mathbf{w}\|_2^2 + \frac{1}{2} \sum_{s=1}^t (\mathbf{w}^\top \mathbf{x}_s - y_s)^2 \quad (24)$$

Letting $A_t = I + \sum_{s=1}^t \mathbf{x}_s \mathbf{x}_s^\top$, where I is the $d \times d$ identity matrix, the Ridge Regression iterates (24) can be written in closed form as $\mathbf{w}_{t+1} = A_t^{-1} (y_1 \mathbf{x}_1 + \dots + y_t \mathbf{x}_t)$. Online Ridge Regression enjoys the following regret bound (Cesa-Bianchi & Lugosi 2006, Theorem 11.7),

$$R_T(\mathbf{u}) \leq \frac{\|\mathbf{u}\|^2}{2} + d \ln \left(1 + \frac{T}{d} \left(\max_{t=1, \dots, T} \|\mathbf{x}_t\|_2^2 \right) \right) \left(\max_{t=1, \dots, T} \ell_t(\mathbf{w}_t) \right) \quad \forall \mathbf{u} \in \mathbb{R}^d$$

This result shows a regret potentially logarithmic in time, except for the extraneous quantity $\max_t \ell_t(\mathbf{w}_t)$ in place of the correct scaling factor $\max_t y_t^2$. If we knew an upper bound Y on $\max_t |y_t|$, then we could obtain the desired scaling by clipping predictions $\mathbf{w}_t^\top \mathbf{x}_t$ in the interval $[-Y, Y]$, (Vovk 2001, Theorem 4). Luckily, there is a better fix which does not require any preliminary knowledge about $\max_t |y_t|$: simply add to the objective function in (24) an extra loss term associated with the data point \mathbf{x}_{t+1} and the value $y_{t+1} = 0$. As we see next, this has the effect of shrinking (towards zero) the linear predictions $\mathbf{w}_t^\top \mathbf{x}_t$, thus adding stability to the algorithm. Cesa-Bianchi & Lugosi (2006) call the resulting algorithm

VAW, after Vovk (2001) and Azoury & Warmuth (2001) who independently introduced it. VAW iterates are defined by

$$\mathbf{w}_{t+1} = \operatorname{argmin}_{\mathbf{w} \in \mathbb{R}^d} \frac{1}{2} \|\mathbf{w}\|_2^2 + \frac{1}{2} \sum_{s=1}^t \left(\mathbf{w}^\top \mathbf{x}_s - y_s \right)^2 + \frac{1}{2} \left(\mathbf{w}^\top \mathbf{x}_{t+1} \right)^2 \quad (25)$$

The closed form expression is simply $\mathbf{w}_{t+1} = A_{t+1}^{-1} (y_1 \mathbf{x}_1 + \dots + y_t \mathbf{x}_t)$. In order to appreciate how predictions are shrunk by the addition of the term $\frac{1}{2} (\mathbf{w}^\top \mathbf{x}_{t+1})^2$, let $\hat{y}_{t+1}^{\text{RR}} = \mathbf{w}_{t+1}^\top \mathbf{x}_{t+1}$ be the Ridge Regression prediction computed via (24). Then, the VAW prediction $\hat{y}_{t+1}^{\text{VAW}} = \mathbf{w}_{t+1}^\top \mathbf{x}_{t+1}$ computed via (25) satisfies

$$\hat{y}_{t+1}^{\text{VAW}} = \frac{\hat{y}_{t+1}^{\text{RR}}}{1 + \mathbf{x}_{t+1}^\top A_t^{-1} \mathbf{x}_{t+1}}.$$

Clearly, $|\hat{y}_{t+1}^{\text{VAW}}| < |\hat{y}_{t+1}^{\text{RR}}|$ whenever $\|\mathbf{x}_{t+1}\|_2 > 0$. It is interesting to compare VAW to a Gaussian Process (GP) for regression (Rasmussen & Williams 2005). It is known that GP predicts with $\hat{y}_{t+1}^{\text{RR}}$, however it also returns an estimate of the variance of the prediction that depends on $\mathbf{x}_{t+1}^\top A_t^{-1} \mathbf{x}_{t+1}$. Hence, VAW shrinks more the points that are assigned a high variance estimate by a GP. The regret analysis of WAV shows that

$$R_T(\mathbf{u}) \leq \frac{\|\mathbf{u}\|^2}{2} + \left(\max_{t=1, \dots, T} y_t^2 \right) \frac{d}{2} \ln \left(1 + \frac{T}{d} \left(\max_{t=1, \dots, T} \|\mathbf{x}_t\|_2^2 \right) \right) \quad \forall \mathbf{u} \in \mathbb{R}^d \quad (26)$$

Note that this bound simultaneously holds for any sequence $(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots \in \mathbb{R}^d \times \mathbb{R}$, for any time horizon T , and for any linear comparator $\mathbf{u} \in \mathbb{R}^d$. Moreover, as proven in Vovk (2001, Section 3.3), the bound (26) is asymptotically optimal, including the leading constant.

3.3. Optimistic updates

The iterates of VAW are computed by adding to the expression in the argmin an extra loss term $\frac{1}{2} (\mathbf{w}^\top \mathbf{x}_{t+1})^2$ that “predicts” the next label y_{t+1} to be zero. This idea can be generalized to what is known in the literature as an *optimistic update* (Chiang et al. 2012, Rakhlin & Sridharan 2013). Given any sequence ℓ_1, ℓ_2, \dots of convex losses, FTRL with linearized losses and optimistic updates uses the iterates

$$\mathbf{w}_{t+1} = \operatorname{argmin}_{\mathbf{w} \in \mathbb{V}} \psi_{t+1}(\mathbf{w}) + \sum_{s=1}^t \mathbf{g}_s^\top \mathbf{w} + \hat{\mathbf{g}}_{t+1}^\top \mathbf{w}$$

where $\mathbf{g}_s = \nabla \ell_s(\mathbf{w}_s)$, $\hat{\mathbf{g}}_{t+1}$ is a guess for $\nabla \ell_{t+1}(\mathbf{w}_{t+1})$, and ψ_1, ψ_2, \dots are arbitrary μ -strongly convex regularizers with respect to a norm $\|\cdot\|$. One can show that, for all $\mathbf{u} \in \mathbb{V}$,

$$R_T(\mathbf{u}) \leq \psi_{T+1}(\mathbf{u}) - \psi_1(\mathbf{w}_1) + \frac{1}{2\mu} \sum_{t=1}^T \left(\|\mathbf{g}_t - \hat{\mathbf{g}}_t\|_* + \psi_t(\mathbf{w}_{t+1}) - \psi_{t+1}(\mathbf{w}_{t+1}) \right) \quad (27)$$

We now describe two concrete and simple examples. First, consider $\hat{\mathbf{g}}_{t+1} = \frac{1}{t} \sum_{s=1}^t \mathbf{g}_s$ with $\hat{\mathbf{g}}_1 = \mathbf{0}$ and—for the sake of simplicity—assume the convex losses are 1-Lipschitz in \mathbb{V} , so that $\|\mathbf{g}_t\|_2 \leq 1$ for all \mathbf{g}_t generated in (27). Take $\psi_t = \frac{\mu}{2} \|\cdot\|_2^2$ for all t . Then, for all $\mathbf{u} \in \mathbb{V}$,

$$R_T(\mathbf{u}) \leq \frac{\mu}{2} \|\mathbf{u}\|_2^2 + \frac{1}{2\mu} \sum_{t=1}^T \|\mathbf{g}_t - \hat{\mathbf{g}}_t\|_2^2$$

Now note that $\ell_t = \frac{1}{2} \|\mathbf{g}_t - \cdot\|_2^2$ is a 1-strongly convex loss, and $\widehat{\mathbf{g}}_1, \widehat{\mathbf{g}}_2, \dots$ correspond to the predictions of FTL on the sequence ℓ_1, ℓ_2, \dots of such losses. Therefore, using the FTL analysis for strongly convex losses mentioned at the beginning of this section we have that

$$\sum_{t=1}^T \|\mathbf{g}_t - \widehat{\mathbf{g}}_t\|_2^2 - \min_{\mathbf{g}: \|\mathbf{g}\|_2 \leq 1} \sum_{t=1}^T \|\mathbf{g}_t - \mathbf{g}\|_2^2 = \mathcal{O}(\ln T)$$

It is easy to see that the minimizer in the above expression is $\bar{\mathbf{g}} = \frac{1}{T} \sum_t \mathbf{g}_t$. So, we conclude

$$R_T(\mathbf{u}) \leq \frac{\mu}{2} \|\mathbf{u}\|_2^2 + \frac{1}{2\mu} \sum_{t=1}^T \|\mathbf{g}_t - \bar{\mathbf{g}}\|_2^2 + \mathcal{O}(\ln T)$$

Hence, the regret is bounded in terms of the cumulative empirical variance of the loss gradients. Using more sophisticated time-dependent regularizers, the regret can be bounded by an expression sublinear in the cumulative variance. We mention a similar result in the next example, where we choose $\widehat{\mathbf{g}}_{t+1} = \mathbf{g}_t$. In other words, we guess the next loss gradient to be similar to the current one. This is only beneficial under the additional assumption that losses are smooth, which is equivalent to assume that gradients are Lipschitz. Then, choosing regularizers $\psi_t = \beta_t \psi$, where ψ is a base 1-strongly convex regularizer with respect to a norm $\|\cdot\|$ and $\beta_t > 0$ is a scaling factor, one can prove that the regret is bounded by an expression of the order of (see Orabona 2019)

$$\sqrt{1 + \sum_{t=2}^T \|\nabla \ell_t(\mathbf{w}_{t-1}) - \nabla \ell_{t-1}(\mathbf{w}_{t-1})\|_*^2}$$

Bounds of this form were first proven by Chiang et al. (2012). Optimistic updates have been also applied to show fast rates for regret minimization problems in game theory by Syrgkanis et al. (2015), Foster et al. (2016).

4. UNCONSTRAINED ONLINE CONVEX OPTIMIZATION

As discussed in previous sections, both Online Mirror Descent (OMD) and Follow the Regularized Leader (FTRL) enjoy a regret of order $\mathcal{O}(D\sqrt{T})$ for convex and Lipschitz losses, where D bounds the diameter of \mathbb{V} according to the divergence D_ψ (for OMD) or to the range of the base regularizer Ψ (for FTRL), see (8) and (14). If \mathbb{V} is unbounded, say $\mathbb{V} = \mathbb{R}^d$, and losses are convex Lipschitz, we can still run OGD with $\mathbf{w}_1 = \mathbf{0}$ and fixed stepsize $\eta = \alpha/\sqrt{T}$ for $\alpha > 0$. Using (10) we get the following upper bound on the regret

$$R_T(\mathbf{u}) \leq \frac{1}{2} \left(\frac{\|\mathbf{u}\|_2^2}{\alpha} + \alpha \right) \sqrt{T} \quad \forall \mathbf{u} \in \mathbb{R}^d \quad (28)$$

By tuning α optimally with respect to $\|\mathbf{u}\|_2^2$ we could get the bound $R_T(\mathbf{u}) \leq \|\mathbf{u}\|_2 \sqrt{T}$. This is equivalent to what we would get by running OGD with projection in the Euclidean ball of radius $U = \|\mathbf{u}\|_2$. However, this bound cannot be simultaneously achieved for all \mathbf{u} because the algorithms must be run with only one choice of their parameters. The problem we study in this section is whether we can get a bound better than (28) in the unconstrained setting. In order to answer this question, we explore the connection between online convex

optimization and the problem of sequentially betting on the elements of a deterministic sequence of values.

The fact that certain instances of online learning could be phrased as a gambling problem is known since the works of Kelly (1956), Cover (1974), and Feder (1991). It was only recently, however, that Orabona & Pál (2016) realized that certain betting strategies could be used to derive OCO algorithms with no parameters to tune, which could be also used for solving unconstrained OCO in an optimal way.

The betting game we are interested in is parameterized by an unknown deterministic sequence $x_1, x_2, \dots \in [-1, 1]$ of real numbers. The bettor starts out with an initial wealth of $C_0 = \varepsilon > 0$. In each round $t = 1, 2, \dots$ of the game,

1. The bettor bets $\alpha_t \in [-1, 1]$, whose absolute value is the fraction of their current wealth C_{t-1} they are betting, and whose sign indicates which sign we bet x_t will have;
2. The next value $x_t \in [-1, 1]$ is revealed;
3. The bettor's return is $x_t \times \alpha_t C_{t-1} \in \mathbb{R}$.

Note that the bettor's wealth changes in each step t according to $C_t = (1 + \alpha_t x_t) C_{t-1}$.

In order to show how to apply a betting strategy to any OCO problems, we first consider the 1-dimensional case with losses ℓ_1, ℓ_2, \dots defined on \mathbb{R} and having uniformly bounded derivatives $|\ell'_t| \leq 1$ for all t . First of all, recall that, because of the convexity of losses, for any $u \in \mathbb{R}$ the regret $R_T(u)$ of any online algorithm predicting with $w_1, w_2, \dots \in \mathbb{R}$ is upper bounded by $\sum_{t=1}^T (w_t - u) g_t$, where g_t is the derivative of ℓ_t at w_t . In order to use a betting strategy to solve the online problem, we set $w_t = \alpha_t C_{t-1}$. We then have $C_t = C_{t-1} + w_t x_t$, which implies $C_T = \varepsilon + \sum_{t=1}^T w_t x_t$. If we now set $x_t = -g_t \in [-1, 1]$ (by our assumption on the losses), we get

$$C_T = \varepsilon \prod_{t=1}^T (1 + \alpha_t x_t) = \varepsilon + \sum_{t=1}^T w_t x_t = \varepsilon - \sum_{t=1}^T w_t g_t \quad \forall x_1, \dots, x_T \quad (29)$$

Next, we prove that a lower bound on the wealth C_T can be used to upper bound $R_T(u)$ for all $u \in \mathbb{R}$. In particular, assume that we have a betting strategy $\alpha_1, \alpha_2, \dots \in \mathbb{R}$ such that

$$C_T \geq \phi \left(\sum_{t=1}^T x_t \right) = \phi \left(- \sum_{t=1}^T g_t \right) \quad (30)$$

for some convex real function ϕ . Then, for any $u \in \mathbb{R}$ we have

$$\begin{aligned}
R_T(u) &\leq \sum_{t=1}^T (w_t - u)g_t \\
&= \sum_{t=1}^T -ug_t - \left(\varepsilon - \sum_{t=1}^T w_t g_t \right) + \varepsilon \\
&= \sum_{t=1}^T -ug_t - C_T + \varepsilon && \text{(using (29))} \\
&\leq \sum_{t=1}^T -ug_t - \phi \left(- \sum_{t=1}^T g_t \right) + \varepsilon && \text{(by our assumption on } C_T) \\
&\leq \sup_{\theta \in \mathbb{R}} \theta u - \phi(\theta) + \varepsilon \\
&= \phi^*(u) + \varepsilon && \text{(using the definition (7) of Fenchel conjugate)}
\end{aligned}$$

Now that we established a formal connection between online prediction (in the 1-dimensional case) and betting, we must pick a good betting strategy, one that has a small regret in the betting game. We measure this regret against the class of strategies that always bet the best constant $\alpha \in [-1, 1]$. After T steps, the final wealth of each such strategy is $C_T(\alpha) = \varepsilon \prod_{t=1}^T (1 + \alpha x_t)$. Since wealth changes multiplicatively, it is natural to define the regret of a betting strategy $\alpha_1, \alpha_2, \dots$ using the logarithm of the wealth,

$$R_T(\alpha) = \ln \left(\varepsilon \prod_{t=1}^T (1 + \alpha x_t) \right) - \ln \left(\varepsilon \prod_{t=1}^T (1 + \alpha_t x_t) \right) = \sum_{t=1}^T \ln(1 + \alpha x_t) - \sum_{t=1}^T \ln(1 + \alpha_t x_t)$$

In the special case of Boolean values $x_t \in \{-1, 1\}$, the *optimal* regret $R_T(\alpha)$ was very precisely determined (including constants) by Shtarkov (1987), who established that

$$R_T(\alpha) = \frac{1}{2} \ln T + \frac{1}{2} \ln \frac{\pi}{2} + o(1) \quad \forall \alpha \in [-1, 1]$$

Unfortunately, Shtarkov's strategy is not efficiently computable. A more efficient (and nearly optimal) strategy was introduced a bit earlier by Krichevsky & Trofimov (1981). The Krichevsky-Trofimov (KT) strategy has a regret of

$$R_T(\alpha) = \frac{1}{2} \ln T + \Theta(1) \quad \forall \alpha \in [-1, 1]$$

which is optimal only up to constants. When $x_t \in [-1, 1]$, the KT strategy takes the simple form $\alpha_1 = 0$ and $\alpha_t = (x_1 + \dots + x_{t-1})/t$. Orabona & Pál (2016) proved that the wealth of the KT strategy satisfies

$$C_T \geq \frac{\varepsilon}{c\sqrt{T}} \exp \left(\frac{1}{2T} \left(\sum_{t=1}^T x_t \right)^2 \right)$$

for some universal constant $c > 0$.

Going back to the 1-dimensional online convex optimization case, recalling (29) and $x_t = -g_t \in [-1, 1]$, we see that the KT strategy generates predictions of the form

$$w_t = \alpha_t C_{t-1} = - \left(\frac{1}{t} \sum_{s=1}^{t-1} g_s \right) \left(\varepsilon - \sum_{s=1}^{t-1} g_s w_s \right) \quad (31)$$

Using (30) and our bound $R_T(u) \leq \phi^*(u) + \varepsilon$ on the regret, we can compute ϕ^* for our case and derive the bound

$$R_T(u) \leq |u| \sqrt{T \ln(u^2 c^2 T + 1)} + 1 \quad \forall u \in \mathbb{R} \quad (32)$$

where, for simplicity, we set $\varepsilon = 1$. Recall that this bound holds, simultaneously for all $u \in \mathbb{R}$, under the condition that $|\nabla \ell_t| \leq 1$ for all t . Under the same conditions, OGD guarantees a bound on the regret $R_T(u)$ of worse order $u^2 \sqrt{T}$, see (28).

This result has been extended to \mathbb{R}^d (and more generally to Banach spaces) by Cutkosky & Orabona (2018) using a simple trick: in order to bound the regret $R_T(\mathbf{u})$ against any $\mathbf{u} \in \mathbb{R}^d$ we use the 1-dimensional KT strategy to learn the length $\|\mathbf{u}\|$ of \mathbf{u} and OMD with model space $\mathbb{V} \equiv \{\mathbf{u} \in \mathbb{R}^d : \|\mathbf{u}\| \leq 1\}$ to learn the direction $\mathbf{u}/\|\mathbf{u}\|$ of \mathbf{u} , where $\|\cdot\|$ is any desired norm. This is done as follows: let w_t the prediction (31) of KT and let \mathbf{v}_t be the prediction of OMD. Then, the combined algorithm predicts $w_t \mathbf{v}_t$ and, upon receiving the gradient $\mathbf{g}_t = \nabla \ell_t(w_t \mathbf{v}_t)$, feeds the derivative $\mathbf{g}_t^\top \mathbf{v}_t$ to KT and the gradient \mathbf{g}_t to OMD.

To analyze the regret of this combined KT-OMD strategy, note that $\|\mathbf{v}_t\| \leq 1$ for all t because OMD projects onto \mathbb{V} . Therefore, $|\mathbf{g}_t^\top \mathbf{v}_t| \leq \|\mathbf{g}_t\|_*$. So, by linearizing the regret, we can write

$$\begin{aligned} R_T(\mathbf{u}) &= \sum_{t=1}^T \ell_t(w_t \mathbf{v}_t) - \sum_{t=1}^T \ell_t(\mathbf{u}) \\ &\leq \sum_{t=1}^T \mathbf{g}_t^\top (w_t \mathbf{v}_t - \mathbf{u}) \\ &= \sum_{t=1}^T \left(w_t \mathbf{g}_t^\top \mathbf{v}_t - \|\mathbf{u}\| \mathbf{g}_t^\top \mathbf{v}_t \right) + \sum_{t=1}^T \left(\|\mathbf{u}\| \mathbf{g}_t^\top \mathbf{v}_t - \mathbf{g}_t^\top \mathbf{u} \right) \\ &= \sum_{t=1}^T \left(w_t \mathbf{g}_t^\top \mathbf{v}_t - \|\mathbf{u}\| \mathbf{g}_t^\top \mathbf{v}_t \right) + \|\mathbf{u}\| \sum_{t=1}^T \left(\mathbf{g}_t^\top \mathbf{v}_t - \mathbf{g}_t^\top \frac{\mathbf{u}}{\|\mathbf{u}\|} \right). \end{aligned}$$

Now note that the first sum is the linearized regret of KT against $\|\mathbf{u}\|$, where $\mathbf{g}_t^\top \mathbf{v}_t$ are the loss derivatives. The second sum (ignoring the $\|\mathbf{u}\|$ factor) is instead the linearized regret of OMD against $\mathbf{u}/\|\mathbf{u}\|$ with loss gradients \mathbf{g}_t .

Since KT requires the loss derivatives to belong to $[-1, 1]$, we can apply the above bound when $|\mathbf{g}_t^\top \mathbf{v}_t| \in [0, 1]$, which holds when ℓ_t is 1-Lipschitz for all t . Assuming that, and in view of comparing to (28), we consider OGD run on the unit Euclidean ball with step size $\eta_t = 1/\sqrt{t}$ and KT run with $\varepsilon = 1$. Then using (9) for OGD and (32) for KT we obtain that the regret of KT-OGD satisfies

$$R_T(\mathbf{u}) = \mathcal{O} \left(\left(\sqrt{\ln(\|\mathbf{u}\|_2^2 T + 1)} + 1 \right) \|\mathbf{u}\|_2 \sqrt{T} + 1 \right) \quad \forall \mathbf{u} \in \mathbb{R}^d \quad (33)$$

Note that, except for the logarithmic factor $\sqrt{\ln(\|\mathbf{u}\|_2^2 T + 1)} + 1$, this bound matches the bound $\|\mathbf{u}\|_2 \sqrt{T}$ for OGD tuned with the unknown knowledge of $\|\mathbf{u}\|$. Indeed, this logarithmic term is the price we have to pay for the adaptivity of the algorithm to $\|\mathbf{u}\|$, as Streeter & McMahan (2012) showed it to be *unavoidable* in the unconstrained setting.

Note that parameter-free algorithms like KT-OMD are useful also when \mathbf{u} lives in a bounded subset $\mathbb{V} \subset \mathbb{R}^d$, but the Bregman divergence D_ψ is unbounded in \mathbb{V} . The prime

example for this scenario is the entropic mirror map ψ on the simplex \mathbb{V} , where D_ψ is the cross-entropy. The regret bound (10) for the corresponding instance of OMD (which we called EG) has the form

$$R_T(\mathbf{u}) = \mathcal{O}\left(\frac{D_\psi(\mathbf{u}, \mathbf{w}_1)}{\eta} + \eta T\right) \quad \forall \mathbf{u} \in \mathbb{V}$$

for any fixed choice of the stepsize η . In the setting of prediction with expert advice (a special case of OCO), an implicitly defined parameter-free version of EG for linear losses called NormalHedge was introduced by Chaudhuri et al. (2009). The bound was later improved by Chernov & Vovk (2010) with another implicitly defined algorithm. Using a different reduction to betting strategies, Orabona & Pál (2016) proved a bound of order $\sqrt{(D_\psi(\mathbf{u}, \mathbf{w}_1) + 1)T}$ for all \mathbf{u} in the simplex using a closed form update. Note that, as before, this bound is equal to the bound which one would get by tuning EG using the (unknown) knowledge of $D_\psi(\mathbf{u}, \mathbf{w}_1)$. Further improvements were obtained by Koolen & Van Erven (2015), who introduced an algorithm called Squint.

One of the nice features of (33) is that $R_T(\mathbf{0})$ does not depend on T . So we can then run multiple instances of the combined KT-OMD strategy using a different mirror map for each OMD instance. This allows to bound the regret as if we ran the algorithm using the best mirror map in hindsight (where best is relative to the linearized regret). Suppose for example we run two instances of KT-OMD and predict using the sum $\mathbf{w}_t = \mathbf{w}_t^{(1)} + \mathbf{w}_t^{(2)}$ of their predictions. Then, linearizing the regret, we obtain

$$\begin{aligned} R_T(\mathbf{u}) &\leq \sum_{t=1}^T \mathbf{g}_t^\top (\mathbf{w}_t - \mathbf{u}) = \min_{\substack{\mathbf{v}, \mathbf{z} \in \mathbb{R}^d \\ \mathbf{v} + \mathbf{z} = \mathbf{u}}} \left(\sum_{t=1}^T \mathbf{g}_t^\top (\mathbf{w}_t^{(1)} - \mathbf{v}) + \sum_{t=1}^T \mathbf{g}_t^\top (\mathbf{w}_t^{(2)} - \mathbf{z}) \right) \\ &= \min \left\{ \sum_{t=1}^T \mathbf{g}_t^\top (\mathbf{w}_t^{(1)} - \mathbf{u}), \sum_{t=1}^T \mathbf{g}_t^\top (\mathbf{w}_t^{(2)} - \mathbf{u}) \right\} \end{aligned}$$

For instance, if we run OMD using the p -norm³ mirror map $\psi = \frac{1}{2} \|\cdot\|_p^2$ for $1 < p \leq 2$, which is $(p-1)$ -strongly convex with respect to the same norm, then $R_T(\mathbf{u})$ is bounded in terms of $\|\mathbf{u}\|_p$ and $\|\mathbf{g}_t\|_q$, where q is the conjugate coefficient of p , satisfying $\frac{1}{p} + \frac{1}{q} = 1$. Interestingly, one can show (Gentile & Littlestone 1999) that choosing $p = 2 \ln d / (2 \ln d - 1)$ (for $d \geq 3$) gives a regret bound of the form $(\|\mathbf{u}\|_1^2 / \alpha + \alpha) G \sqrt{T \ln d}$ for all $\mathbf{u} \in \mathbb{R}^d$ and for $\max_t \|\mathbf{g}_t\|_\infty \leq G$. As this is very similar to the bound obtained by EG when \mathbb{V} is the simplex, we see that by choosing $p \in (1, 2]$ we can interpolate between OGD and EG. By running multiple instances of KT-OMD, where the OMD instances use different values of p , one can derive a regret bound almost as good as the p -norm OMD run with the best value of p .

5. OTHER NOTIONS OF REGRET

Bounding the regret $R_T(\mathbf{u})$ for all $\mathbf{u} \in \mathbb{V}$ may not be crucial in some practical applications. For example, if the loss sequence ℓ_1, ℓ_2, \dots is such that no $\mathbf{u} \in \mathbb{V}$ achieves a small cumulative loss $\ell_1(\mathbf{u}) + \ell_2(\mathbf{u}) + \dots$, then regret bounds may not be at all helpful in telling good algorithms from bad ones. This lack of a single good minimizer in \mathbb{V} of the cumulative loss is

³Recall the definition of p -norm of a vector $\mathbf{u} \in \mathbb{R}^d$, $\|\mathbf{u}\|_p = (|u_1|^p + \dots + |u_d|^p)^{1/p}$.

likely to occur when the loss sequence is generated by a highly nonstationary data sequence, possibly affected by seasonalities and other disturbances. In this case, regret should be replaced by more robust measures, allowing better comparators than fixed elements of \mathbb{V} . In what follows, D_2 is the Euclidean diameter of \mathbb{V} and L is the Lipschitz constant of the convex loss functions in the sequence ℓ_1, \dots, ℓ_T .

Dynamic regret. A notion of regret which captures nonstationary comparators is that of *dynamic regret* (Herbster & Warmuth 1998b),

$$R_T^{\text{dyn}}(\mathbf{u}_1, \dots, \mathbf{u}_T) = \sum_{t=1}^T \ell_t(\mathbf{w}_t) - \sum_{t=1}^T \ell_t(\mathbf{u}_t) \quad \text{where } \mathbf{u}_1, \dots, \mathbf{u}_T \in \mathbb{V}$$

Note that $R_T^{\text{dyn}}(\mathbf{u}, \dots, \mathbf{u}) = R_T(\mathbf{u})$, so dynamic regret includes standard regret as a special case. Zhang et al. (2018a) show a general lower bound on dynamic regret of the form $\Omega\left(L\sqrt{(D_2 + \Pi_{2,T})D_2T}\right)$, where

$$\Pi_{p,T} = \sum_{t=1}^{T-1} \|\mathbf{u}_{t+1} - \mathbf{u}_t\|_p$$

is the p -norm path-length function, measuring the nonstationarity of the comparator sequence $\mathbf{u}_1, \dots, \mathbf{u}_T$. When $\mathbf{u}_1 = \dots = \mathbf{u}_T$ then $\Pi_{p,T} = 0$. In this case the lower bound on the dynamic regret reduces (for $p = 2$) to the lower bound $\Omega\left(L_2 D_2 \sqrt{T}\right)$ on the standard regret proven in Section 1.3.

Herbster & Warmuth (1998b)—see also Cesa-Bianchi & Lugosi (2006, Theorem 11.4)—prove upper bounds on the dynamic regret of OMD run with the p -norm mirror map (for $1 < p \leq 2$) of the form

$$R_T^{\text{dyn}}(\mathbf{u}_1, \dots, \mathbf{u}_T) \leq \frac{\|\mathbf{u}_T\|_p^2 + \Pi_{p,T} D_p}{\eta} + \eta L_q^2 T \quad \forall \mathbf{u}_1, \dots, \mathbf{u}_T \in \mathbb{V}$$

where D_p is the diameter of \mathbb{V} measured using the p -norm, and L_q is the Lipschitz constant (with respect to the dual q -norm) of the loss functions in the sequence ℓ_1, \dots, ℓ_T —see also Zinkevich (2003) for similar results in the special case of $p = 2$. Choosing the stepsize as $\eta = 1/(L_q \sqrt{T})$ gives a suboptimal upper bound of the order of $L_q(D_p^2 + \Pi_{p,T} D_p) \sqrt{T}$. In the Euclidean case ($p = 2$), Zhang et al. (2018a) use the Hedge algorithm (a special case of EG) to aggregate $\mathcal{O}(\ln T)$ instances of OGD, each one run with a different choice of η to guess the desired value of $\Pi_{2,T}$ (up to a constant factor). They prove the dynamic regret bound $\mathcal{O}\left(L_2 \sqrt{(D_2 + \Pi_{2,T})D_2T}\right)$, matching the lower bound up to constants.

Adaptive regret. A different view on the theme of nonstationary comparators is offered by the notion of *adaptive regret* (Hazan & Seshadhri 2007). Adaptive regret evaluates the performance of the online algorithm against that of the best fixed comparator in any interval of time. Formally,

$$R_{\tau,T}^{\text{ada}} = \max_{s=1, \dots, T-\tau+1} \left(\sum_{t=s}^{s+\tau-1} \ell_t(\mathbf{w}_t) - \min_{\mathbf{u} \in \mathbb{V}} \sum_{t=s}^{s+\tau-1} \ell_t(\mathbf{u}) \right), \quad \text{where } \tau \in \{1, \dots, T\}$$

In their paper, Hazan & Seshadhri (2007) use a harder notion of adaptive regret, namely $\max_{\tau} R_{\tau,T}^{\text{ada}}$. They show an online algorithm whose regret grows in T like $\sqrt{T(\ln T)^3}$. In a

follow-up paper, Daniely et al. (2015) devise an online algorithm with the adaptive regret bound

$$R_{\tau,T}^{\text{ada}}(\mathbf{u}) = \mathcal{O}\left((DL + \ln T) \sqrt{\tau}\right) \quad \mathbf{u} \in \mathbb{V} \quad (34)$$

where D is the Euclidean diameter of \mathbb{V} and L is the Lipschitz constant (with respect to the Euclidean norm) of the loss functions in the sequence ℓ_1, \dots, ℓ_T . This result is improved by Jun et al. (2017), who show the better bound $R_{\tau,T}^{\text{ada}}(\mathbf{u}) = \mathcal{O}\left(\left(DL + \sqrt{\ln T}\right) \sqrt{\tau}\right)$ using the betting framework described in Section 4.

Most of the online algorithms for minimizing adaptive regret work by combining several instances of an online algorithm for the standard notion of regret. Each instance is run in a specific interval of time, where the set of intervals is carefully designed so that the overall number of instances to be run is $\mathcal{O}(\ln T)$. These instances are then combined using an algorithm based on the framework of prediction with expert advice (Cesa-Bianchi et al. 1997) where each instance is viewed as an expert. As instances typically run for less than T time steps, Jun et al. (2017) combine the betting framework with the sleeping experts model (Freund et al. 1997), which allows only for a subset of the experts to be active at any point of time.

Although the algorithm of Jun et al. (2017) is designed to minimize adaptive regret, Zhang et al. (2018b) show that the same algorithm can be also used to prove the following dynamic regret result

$$R_T^{\text{dyn}}(\mathbf{u}_1, \dots, \mathbf{u}_T) = \mathcal{O}\left(DL \max\left\{\sqrt{T \ln T}, T^{2/3} V_T^{1/3} (\ln T)^{1/3}\right\}\right) \quad (35)$$

where

$$V_T = \sum_{t=2}^T \sup_{\mathbf{u} \in \mathbb{V}} \|\ell_t(\mathbf{u}) - \ell_{t-1}(\mathbf{u})\|_2$$

measures the variation of the loss sequence ℓ_1, \dots, ℓ_T . As shown by Besbes et al. (2015), the dependence on V_T in (35) is not improvable.

A result relating dynamic regret to adaptive regret is proven by Zhang et al. (2018b), who show that

$$R_T^{\text{dyn}}(\mathbf{u}_1, \dots, \mathbf{u}_T) \leq \min_{\mathcal{P}_T} \sum_{\mathcal{I} \in \mathcal{P}_T} \left(R_{|\mathcal{I}|,T}^{\text{ada}} + 2|\mathcal{I}|V_T(\mathcal{I})\right)$$

where the min is taken over all partitions \mathcal{P}_T of $\{1, \dots, T\}$ in intervals $\mathcal{I} = \{t_r, \dots, t_s\}$ of consecutive time steps, with $1 \leq r \leq s \leq T$. The quantity

$$V_T(\mathcal{I}) = \sum_{t \in \mathcal{I}} \sup_{\mathbf{u} \in \mathbb{V}} \|\ell_t(\mathbf{u}) - \ell_{t-1}(\mathbf{u})\|_2$$

is the variation of the loss sequence within the time interval. This is later extended by Zhang et al. (2020), who prove

$$R_T^{\text{dyn}}(\mathbf{u}_1, \dots, \mathbf{u}_T) \leq \min_{\mathcal{P}_T} \sum_{\mathcal{I} \in \mathcal{P}_T} \left(R_{|\mathcal{I}|,T}^{\text{ada}} + L|\mathcal{I}|\Pi_T(\mathcal{I})\right) \quad \forall \mathbf{u}_1, \dots, \mathbf{u}_T \in \mathbb{V}$$

where

$$\Pi_T(\mathcal{I}) = \sum_{t \in \mathcal{I}} \|\mathbf{u}_{t+1} - \mathbf{u}_t\|_2$$

is the path length over the interval \mathcal{I} . However, when combined with known bounds on the dynamic regret, this bound does not give the optimal bound $\mathcal{O}\left(L_2 \sqrt{(D_2 + \Pi_{2,T})D_2 T}\right)$

for dynamic regret. Zhang et al. (2020) also derive algorithms simultaneously minimizing adaptive and dynamic regret.

Some of these notions of regret were originally introduced in the setting of prediction with expert advice (i.e., OCO with linear losses, where \mathbb{V} is equal to the probability simplex and regret is measured against the corners of the simplex, where linear functions are minimized). In that framework, dynamic regret is known as *tracking* or *shifting* regret (Herbster & Warmuth 1998a). Other notions of regret, instead, are mostly studied in the experts framework. For example, policy regret (Arora et al. 2012a) applies to settings where the loss function ℓ_t at each time t depends not only on the current model \mathbf{w}_t , but also on the past models \mathbf{w}_{t-s} , where s spans a window in the past (whose size H could potentially depend on t). These loss functions can be used to model natural scenarios, such as the “switching cost” scenario where $H = 1$ and $\ell_t(\mathbf{w}_t, \mathbf{w}_{t+1}) = c$ whenever $\mathbf{w}_t \neq \mathbf{w}_{t-1}$ (Kalai & Vempala 2005). Swap regret (Blum & Mansour 2007), instead, measures regret against a set of modification rules. Each modification rule F is an operator on the set $\{1, \dots, d\}$ of coordinates. The instantaneous regret at time t against F of an algorithm choosing \mathbf{w} in the simplex is $\ell_t(\mathbf{w}) - \ell_t(\mathbf{w}^{(F)})$, where $w_i^{(F)} = \sum_{j: F(j)=i} w_j$. Note that when $F(j) = i$ for all $j = 1, \dots, d$, swap regret against F reduces to standard regret against the i -th corner \mathbf{e}_i of the simplex, that is $\ell_t(\mathbf{w}^{(F)}) = \ell_t(\mathbf{e}_i)$ for all \mathbf{w} . Swap regret is especially important when using online learning algorithms to approximate equilibria in games (Cesa-Bianchi & Lugosi 2006).

DISCLOSURE STATEMENT

The authors are not aware of any affiliations, memberships, funding, or financial holdings that might be perceived as affecting the objectivity of this review.

ACKNOWLEDGMENTS

NCB is partially supported by the MIUR PRIN grant *Algorithms, Games, and Digital Markets* (ALGADIMAR). FO is partially supported by the National Science Foundation under grant no. 1925930 *Collaborative Research: TRIPODS Institute for Optimization and Learning*.

LITERATURE CITED

- Abernethy J, Hazan E, Rakhlin A. 2008. Competing in the dark: An efficient algorithm for bandit linear optimization, In *Proceedings of the 21st Annual Conference on Learning Theory*
- Agarwal N, Gonen A, Hazan E. 2019. Learning in non-convex games with an optimization oracle, In *Proceedings of the 32nd Annual Conference on Learning Theory*
- Arora R, Dekel O, Tewari A. 2012a. Online bandit learning against an adaptive adversary: from regret to policy regret, In *Proceedings of the 29th International Conference on Machine Learning*
- Arora S, Hazan E, Kale S. 2012b. The multiplicative weights update method: a meta-algorithm and applications. *Theory of Computing* 8:121–164
- Azoury KS, Warmuth MK. 2001. Relative loss bounds for on-line density estimation with the exponential family of distributions. *Machine Learning* 43:211–246
- Beck A, Teboulle M. 2003. Mirror descent and nonlinear projected subgradient methods for convex optimization. *Operations Research Letters* 31:167–175

- Berger JO. 2013. Statistical decision theory and bayesian analysis. Springer Science & Business Media
- Besbes O, Gur Y, Zeevi A. 2015. Nonstationary stochastic optimization. *Operations research* 63:1227–1244
- Blackwell D. 1956. An analog of the minimax theorem for vector payoffs. *Pacific Journal of Mathematics* 6:1–8
- Block HD. 1962. The Perceptron: A model for brain functioning. *Reviews of Modern Physics* 34:123
- Blum A, Mansour Y. 2007. From external to internal regret. *Journal of Machine Learning Research* 8:1307–1324
- Borodin A, El-Yaniv R. 2005. Online computation and competitive analysis. Cambridge University Press
- Bubeck S. 2015. Convex optimization: Algorithms and complexity. *Foundations and Trends in Machine Learning* 8:231–357
- Cesa-Bianchi N, Freund Y, Haussler D, Helmbold DP, Schapire RE, Warmuth MK. 1997. How to use expert advice. *Journal of the ACM* 44:427–485
- Cesa-Bianchi N, Lugosi G. 2006. Prediction, learning, and games. Cambridge University Press
- Chaudhuri K, Freund Y, Hsu DJ. 2009. A parameter-free hedging algorithm, In *Advances in Neural Information Processing Systems 22*
- Chernov A, Vovk V. 2010. Prediction with advice of unknown number of experts, In *Proceedings of the 26th Conference on Uncertainty in Artificial Intelligence*
- Chiang CK, Yang T, Lee CJ, Mahdavi M, Lu CJ, et al. 2012. Online optimization with gradual variations, In *Proceedings of the 25th Annual Conference on Learning Theory*
- Cortes C, Vapnik V. 1995. Support-vector networks. *Machine Learning* 20:273–297
- Cover T. 1965. Behaviour of sequential predictors of binary sequences, In *Proceedings of the 4th Prague Conference on Information Theory, Statistical Decision Functions and Random Processes*. Publishing House of the Czechoslovak Academy of Sciences
- Cover TM. 1974. Universal gambling schemes and the complexity measures of Kolmogorov and Chaitin. *Technical Report 12*
- Cutkosky A, Orabona F. 2018. Black-box reductions for parameter-free online learning in Banach spaces, In *Proceedings of the 31st Annual Conference on Learning Theory*, vol. 75 of *Proceedings of Machine Learning Research*. PMLR
- Dalalyan A, Tsybakov AB. 2008. Aggregation by exponential weighting, sharp PAC-Bayesian bounds and sparsity. *Machine Learning* 72:39–61
- Dalalyan AS, Salmon J, et al. 2012. Sharp oracle inequalities for aggregation of affine estimators. *The Annals of Statistics* 40:2327–2355
- Daniely A, Gonen A, Shalev-Shwartz S. 2015. Strongly adaptive online learning, In *Proceedings of the 32nd International Conference on Machine Learning*
- De Rosa R, Orabona F, Cesa-Bianchi N. 2015. The ABACOC algorithm: a novel approach for nonparametric classification of data streams, In *2015 IEEE International Conference on Data Mining*. IEEE
- Duchi J, Hazan E, Singer Y. 2011. Adaptive subgradient methods for online learning and stochastic optimization. *Journal of Machine Learning Research* 12:2121–2159
- Feder M. 1991. Gambling using a finite state machine. *IEEE Transactions on Information Theory* 37:1459–1465
- Feder M, Merhav N, Gutman M. 1992. Universal prediction of individual sequences. *IEEE Transactions on Information Theory* 38:1258–1270
- Foster DJ, Li Z, Lykouris T, Sridharan K, Tardos E. 2016. Learning in games: Robustness of fast convergence, In *Advances in Neural Information Processing Systems 29*
- Freund Y, Schapire RE. 1997. A decision-theoretic generalization of on-line learning and an application to boosting. *Journal of Computer and System Sciences* 55:119–139
- Freund Y, Schapire RE. 1999. Large margin classification using the Perceptron algorithm. *Machine*

Learning 37:277–296

- Freund Y, Schapire RE, Singer Y, Warmuth MK. 1997. Using and combining predictors that specialize, In *Proceedings of the 29th Annual ACM Symposium on Theory of Computing*
- Fudenberg D, Levine DK. 1995. Consistency and cautious fictitious play. *Journal of Economic Dynamics and Control* 19:1065–1089
- Gentile C, Littlestone N. 1999. The robustness of the p -norm algorithms, In *Proceedings of the 12th Annual Conference on Computational Learning Theory*
- Hannan J. 1957. Approximation to Bayes risk in repeated play. *Contributions to the Theory of Games* 3:97–139
- Hazan E. 2016. Introduction to online convex optimization. *Foundations and Trends in Optimization* 2:157–325
- Hazan E, Agarwal A, Kale S. 2007. Logarithmic regret algorithms for online convex optimization. *Machine Learning* 69:169–192
- Hazan E, Megiddo N. 2007. Online learning with prior knowledge, In *International Conference on Computational Learning Theory*. Springer
- Hazan E, Seshadhri C. 2007. Adaptive algorithms for online decision problems. *Electronic Colloquium on Computational Complexity* 14
- Herbster M, Warmuth MK. 1998a. Tracking the best expert. *Machine learning* 32:151–178
- Herbster M, Warmuth MK. 1998b. Tracking the best regressor, In *Proceedings of the 11th Annual Conference on Computational Learning Theory*
- Hoerl AE, Kennard RW. 2000. Ridge regression: biased estimation for nonorthogonal problems. *Technometrics* 42:80–86
- Jun KS, Orabona F, Wright S, Willett R. 2017. Improved strongly adaptive online learning using coin betting, In *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics*
- Kalai A, Vempala S. 2005. Efficient algorithms for online decision problems. *Journal of Computer and System Sciences* 71:291–307
- Kelly J. 1956. A new interpretation of information rate. *The Bell System Technical Journal* 35:917–926
- Kivinen J, Warmuth MK. 1997. Exponentiated gradient versus gradient descent for linear predictors. *Information and Computation* 132:1–63
- Kivinen J, Warmuth MK. 1999. Averaging expert predictions, In *Proceedings of the European Conference on Computational Learning Theory*. Springer
- Koolen WM, Van Erven T. 2015. Second-order quantile methods for experts and combinatorial games, In *Proceedings of the 28th Annual Conference on Learning Theory*
- Krichevsky R, Trofimov V. 1981. The performance of universal encoding. *IEEE Transactions on Information Theory* 27:199–207
- Kulis B, Bartlett PL. 2010. Implicit online learning, In *Proceedings of the 27th International Conference on Machine Learning*
- Kuzborskij I, Cesa-Bianchi N. 2017. Nonparametric online regression while learning the metric, In *Advances in Neural Information Processing Systems*
- Littlestone N, Warmuth MK. 1994. The weighted majority algorithm. *Information and Computation* 108:212–261
- Luo H, Agarwal A, Cesa-Bianchi N, Langford J. 2016. Efficient second order online learning by sketching, In *Advances in Neural Information Processing Systems* 29
- McMahan HB. 2017. A survey of algorithms and analysis for adaptive online learning. *Journal of Machine Learning Research* 18:1–50
- McMahan HB, Streeter MJ. 2010. Adaptive bound optimization for online convex optimization, In *Proceedings of the 23rd Annual Conference on Learning Theory*. Omnipress
- Mokhtari A, Ozdaglar A, Pattathil S. 2019. A unified analysis of extra-gradient and optimistic gradient methods for saddle point problems: Proximal point approach. *ArXiv preprint*

arXiv:1901.08511

- Nemirovsky AS, Yudin DB. 1983. Problem complexity and method efficiency in optimization. J. Wiley and Sons
- Novikoff AB. 1963. On convergence proofs for Perceptrons. Tech. rep., SRI, Menlo Park, CA
- Orabona F. 2019. A modern introduction to online learning. *arXiv preprint arXiv:1912.13213*
- Orabona F, Pál D. 2016. Coin betting and parameter-free online learning, In *Advances in Neural Information Processing Systems 29*
- Orabona F, Pál D. 2018. Scale-free online learning. *Theoretical Computer Science* 716:50–69. Special Issue on ALT 2015
- Puterman ML. 2014. Markov decision processes: discrete stochastic dynamic programming. John Wiley & Sons
- Rakhlin A, Sridharan K. 2013. Online learning with predictable sequences, In *Proceedings of the 26th Annual Conference on Learning Theory*
- Rakhlin A, Sridharan K. 2014. Statistical learning and sequential prediction. Book draft
- Rasmussen CE, Williams CKI. 2005. Gaussian processes for machine learning. The MIT Press
- Rigollet P, Tsybakov AB, et al. 2012. Sparse estimation by exponential weighting. *Statistical Science* 27:558–575
- Robbins H. 1951. Asymptotically subminimax solutions of compound statistical decision problems, In *Proceedings of the 2nd Berkeley Symposium on Mathematical Statistics and Probability*. University of California Press
- Rosenblatt F. 1958. The Perceptron: a probabilistic model for information storage and organization in the brain. *Psychological Review* 65:386
- Shalev-Shwartz S. 2007. Online learning: Theory, algorithms, and applications. Ph.D. thesis, The Hebrew University of Jerusalem
- Shalev-Shwartz S. 2012. Online learning and online convex optimization. *Foundations and Trends in Machine Learning* 4:107–194
- Shalev-Shwartz S, Singer Y, Srebro N, Cotter A. 2011. Pegasos: Primal estimated sub-gradient solver for SVM. *Mathematical Programming* 127:3–30
- Shtarkov YM. 1987. Universal sequential coding of single messages. *Problemy Peredachi Informatsii* 23:3–17
- Streeter M, McMahan B. 2012. No-regret algorithms for unconstrained online convex optimization. In *Advances in Neural Information Processing Systems 25*. 2402–2410
- Syrkanis V, Agarwal A, Luo H, Schapire RE. 2015. Fast convergence of regularized learning in games, In *Advances in Neural Information Processing Systems 28*
- Vovk V. 2001. Competitive on-line statistics. *International Statistical Review* 69:213–248
- Zhang L, Lu S, Yang T. 2020. Minimizing dynamic regret and adaptive regret simultaneously. *arXiv preprint arXiv:2002.02085*
- Zhang L, Lu S, Zhou ZH. 2018a. Adaptive online learning in dynamic environments, In *Advances in Neural Information Processing Systems 31*
- Zhang L, Yang T, Zhou ZH, et al. 2018b. Dynamic regret of strongly adaptive methods, In *Proceedings of the 35th International Conference on Machine Learning*
- Zinkevich M. 2003. Online convex programming and generalized infinitesimal gradient ascent, In *Proceedings of the 20th International Conference on Machine Learning*
- Zinkevich M. 2004. Theoretical guarantees for algorithms in multi-agent settings. Ph.D. thesis, School of Computer Science, Carnegie Mellon University